



## 0. Executive Summary

**Inference Engineering Today – Top 10 Themes:** In the last 12–18 months, large-language model (LLM) **inference at scale** has confronted critical technical bottlenecks – and spawned new solutions – across every layer of the stack. Below are ten key themes shaping the state-of-the-art, each with its impact on latency, throughput, cost, and reliability, and an assessment of business potential:

1. **Long Context & Memory Management:** *LLMs' memory usage scales with sequence length; new techniques like PagedAttention treat the attention key-value cache like virtual memory to avoid waste* [1](#) [2](#). **Why it matters:** Enables 8k–32k token contexts without running out of GPU memory or crashing latency SLAs, unlocking long-document and chat history use-cases. **Business potential:** **High** – Organizations need long-context models for enterprise documents and assistants, but struggle with *memory costs and crashes*; solutions that efficiently handle long contexts can save significant GPU spend and unlock new applications (strong demand for hosted services or appliances).
2. **Continuous Batching & Smart Scheduling:** *Traditional request-by-request processing leaves GPUs underutilized; "continuous batching" and iteration-level schedulers fill every GPU cycle by serving multiple requests token-by-token in parallel* [3](#) [4](#). **Why it matters:** This **dramatically boosts throughput** (up to 23x in research tests [5](#)) while keeping latency low – crucial for high QPS services and multi-user platforms. **Business potential:** **High** – Every LLM API provider wants to maximize tokens/sec per dollar; a drop-in library or cloud service that optimizes batching and scheduling (like vLLM or improved orchestration) is immediately valuable [6](#) [7](#).
3. **Optimized KV Cache Design:** *The attention KV cache – storing past token states – can exceed model weights in memory footprint (e.g. 1.7 GB per sequence for LLaMA-13B* [8](#)*). Leading engines now share and recycle cache memory across prompts, evict promptly at end-of-request, and even share identical prefixes across requests* [9](#) [10](#). **Why it matters:** Eliminating cache fragmentation and redundancy yields near **zero waste** (under 4% overhead vs 60–80% in older systems [2](#) [11](#)), allowing **more concurrent sessions** per GPU and avoiding OOM failures. **Business potential:** **Medium** – Best-in-class open source (vLLM, NVIDIA's frameworks) already implement this; a commercial opportunity may lie in integrating these techniques into user-friendly enterprise platforms or multi-tenant cluster managers.
4. **Quantization & Low-Precision Inference:** *Techniques like GPTQ and AWQ compress model weights to 4-bit or 8-bit with minimal accuracy loss* [12](#) [13](#). *Coupled with hardware support (NVIDIA's FP8, INT8 on AWS Inferentia, etc.), these enable smaller memory footprint and faster computation.* **Why it matters:** Quantization can **3–4x boost throughput** and cut costs [14](#), often with <1% accuracy drop – a game-changer for deploying 30B+ models under tight budgets. **Business potential:** **High** – Many companies lack in-house expertise to safely quantize or optimize models. Services that automate quantization, optimize models for specific hardware, or offer quantized-ready model hubs can command strong demand (as evidenced by burgeoning tooling around GPTQ and NVIDIA's TensorRT support for 4-bit).
5. **High-Efficiency Model Architectures:** *Emerging model designs (e.g. Grouped-Query Attention in LLaMA-2, multi-query attention in PaLM, or state-space models like Mamba* [15](#)*) aim to maintain quality with less computation.* **Why it matters:** These approaches reduce inference cost per token

– for example, Mamba achieves **5x throughput vs Transformers** by replacing attention with a linear-time recurrent mechanism <sup>15</sup> (Research-grade). Such models can handle ultra-long sequences (hundreds of thousands of tokens) which standard transformers cannot. **Business potential: Medium (Early)** – Still **research-grade**; not yet mainstream. However, startups that commercialize *attention-alternatives* (for niche cases like million-token processing or on-device models) could gain an edge if/when industry begins adopting these architectures.

**6. Decoding Acceleration (Speculative & Parallel Decoding):** *New algorithms use a “draft-and-refine” approach: a smaller model generates several tokens ahead which the large model then quickly verifies, effectively skipping work* <sup>16</sup> <sup>17</sup>. Other strategies attempt multi-token parallel generation. **Why it matters:** Speculative decoding can yield **2-3x faster generation** without model changes <sup>17</sup> – directly improving end-user experience (snappier responses) and cutting cloud costs. **Business potential: Medium** – This is *rapidly moving from research to practice* (OpenAI, vLLM, and Hugging Face have begun integrating it <sup>16</sup> <sup>17</sup>). There’s room for products that provide easy libraries or services to apply speculative decoding with any model, though big players may bake this in, limiting standalone market size.

**7. Heterogeneous & Distributed Serving (Multi-GPU, Multi-Node):** *To serve larger 70B+ models or to increase throughput, inference often spans multiple GPUs (model sharding, pipeline parallelism) and even multiple nodes (clusters with high-speed interconnects)* <sup>18</sup> <sup>19</sup>. **Why it matters:** Efficiently partitioning models and coordinating distributed inference is hard: it requires managing GPU communication (NVLink/InfiniBand), syncing generation steps, and avoiding stragglers. Solutions like **NVIDIA NeMo/TensorRT-LLM with pipeline parallelism** and Run:AI’s “gang scheduling” for multi-node jobs <sup>20</sup> <sup>21</sup> tackle these issues. **Business potential: Medium** – Enterprises with giant models (or needing *high availability across a cluster*) will invest in robust distributed serving. Specialized orchestrators or hosted solutions that simplify multi-GPU inference (especially across commodity hardware) can find a market, though competition from NVIDIA and cloud providers is intense.

**8. Multi-Tenancy, Isolation, and Fairness:** *Inference servers increasingly serve many clients or models on shared GPUs. The challenge is to isolate each tenant’s workload (both for data safety and quality-of-service): no one user’s large request should stall others* <sup>22</sup> <sup>23</sup>. **Why it matters:** In enterprise settings and APIs, you need **consistent tail latency (P95/P99)** and to prevent any one user from consuming all throughput. Techniques like **iteration-level scheduling** (serving all requests in small increments) ensure fairness <sup>24</sup> <sup>25</sup>, and per-tenant quotas or cache partitions prevent cross-talk. **Business potential: High** – This is a *pain point for platforms*: ad-hoc fixes are common (manual priority settings, separate model instances per tenant). A product that offers strong multi-tenant scheduling, fairness policies, and safety guardrails (e.g. preventing prompt data leaks between users) would be highly appealing to cloud providers and SaaS platforms.

**9. Observability & Cost Analytics for LLMs:** *Traditional APM tools don’t understand “tokens” or why an LLM response was slow. New approaches instrument token-level metrics – e.g. tokens per second, per-request GPU time* <sup>26</sup> <sup>27</sup>, *prompt vs generation latency – and map these to cost*. **Why it matters:** Teams need to **track \$/token and latency per prompt** to manage budgets and SLA commitments <sup>28</sup> <sup>29</sup>. For example, spotting a single user’s queries causing disproportionate load, or measuring how a new model version affects tail latency, requires specialized tracing. **Business potential: High** – Many companies will pay for better **LLM observability** (several startups and open-source projects are emerging here <sup>30</sup> <sup>29</sup>). A tool that plugs into any LLM serving stack and provides dashboards for per-tenant usage, latency breakdowns, and anomaly detection could become indispensable in the era of expensive model deployments.

**10. Reliability & Safety in Production Inference:** With LLMs powering products, uptime and safety are as important as performance. This spans autoscaling reliably, hot-swapping models without downtime, graceful handling of OOM or timeouts, and enforcing content/policy filters at inference. **Why it matters:** Users expect the AI to be **always available and behave well** – a single outage or a single toxic output can be costly. Engineering techniques include redundant model replicas across zones, using CPU fallback if GPUs are saturated, and injecting guardrail checks in the generation loop (for content filtering). **Business potential: Medium** – Much of this is *platform engineering* that larger companies build in-house. However, a focused solution (e.g. a managed inference platform emphasizing high availability, or an API add-on for real-time safety filtering) can attract those who cannot afford deep in-house ML infra teams.

Each theme above addresses core **latency-throughput-cost trade-offs** or **production constraints**. The frontier is defined by techniques to push more tokens through per second, handle longer inputs, and serve more users – all **without** exploding cost or sacrificing reliability. From a startup perspective, **key opportunities** cluster around providing these cutting-edge capabilities “as a service” or in easier workflows for the many teams now struggling to productionize LLMs. In the sections below, we map the technical landscape in layers (hardware to workload), dive deep into the toughest problems and state-of-the-art solutions, and identify concrete gaps where new products or services could thrive.

---

## 1. Layered Landscape Map

Modern LLM inference can be visualized as a **layered stack**, from metal to user-facing applications. Each layer has its own constraints and bottlenecks, and friction often occurs at the interfaces between layers. We describe each layer, its challenges, and where the hardest engineering trade-offs lie:

### 1.1 Hardware Layer (GPUs, Memory, Interconnects)

**Key Hardware:** Today’s LLMs mostly run on NVIDIA GPUs (A10, A100, H100, L4/L40, etc.), with growing interest in alternatives (AWS Inferentia/Trainium, AMD MI250/300, etc.). GPUs offer massive parallelism and high-bandwidth memory (HBM) critical for transformer workloads.

- **Memory (VRAM) Constraints:** LLMs are memory-hungry. A 7B model can use ~15GB GPU memory in 16-bit; a 70B can exceed 70GB – **larger than one GPU**. Moreover, **KV cache** memory grows with sequence length; a 13B model at 2048 tokens uses ~1.7GB for cache per sequence<sup>8</sup>. GPUs have limited HBM (e.g. 40GB on A100, 80GB on H100) which becomes a hard cap on model size or batch size. This forces trade-offs: use model parallelism (split model across GPUs) or quantize to fit memory. **HBM Bandwidth** is also finite – ~1.5–2 TB/s on A100/H100 – and long contexts can saturate memory bandwidth moving Q,K,V data around<sup>31 32</sup>. **Bottleneck:** Running very large models or long sequences often becomes *memory-bound*, not compute-bound, meaning the GPU is waiting on data movement.
- **Compute Throughput:** GPUs excel at matrix multiplies (the core of transformer layers). Newer GPUs (H100) have *Tensor Cores* with FP8/FP16 support that can dramatically speed up if models use those precisions. Inference is typically less compute-intensive than training (no backprop), but still, generating each token requires a forward pass through all layers. Thus latency roughly scales with model depth – e.g. a 70B parameter model may take ~200–300 ms for one token on one A100 (depending on optimizations). *Compute vs Memory trade-off:* On some inputs, the GPU may not fully utilize compute because it’s memory-bound (especially for long context attention).

**Bottleneck:** For smaller models or short contexts, **compute** can be the limiter (GPU at 100% utilization); for long contexts or many concurrent sequences, **memory/bandwidth** is the limiter (GPU compute under-utilized while swapping data).

- **Interconnects (GPU-GPU, GPU-CPU):** When models or workloads exceed one GPU, interconnect speed becomes crucial. **Within a server**, NVIDIA's NVLink/NVSwitch provides ~600–900 GB/s links between GPUs – high enough to split model layers across GPUs with tolerable overhead. **Across servers**, InfiniBand or Ethernet (10–200 Gbps, i.e. 1.25–25 GB/s) is *orders of magnitude slower*, so multi-node inference demands careful partitioning to minimize cross-node communication. **Bottleneck:** Multi-node deployment can easily bottleneck on network if each token requires sending activations or KV data between machines. This is why many <400B models stick to a single node with 8–16 GPUs if possible, to leverage fast NVSwitch (e.g. NVIDIA DGX systems)<sup>18</sup>. The hardware trade-off: using more but smaller GPUs vs fewer but bigger GPUs. (E.g. 2xA100 40GB vs 1xH100 80GB – the latter avoids any multi-GPU overhead but might cost more).
- **Multi-Node vs Single-Node:** *Where the hardest trade-off lies:* If a model must span nodes (e.g. a 400B model needing >8×80GB = 640GB total, beyond one node's capacity), one must accept **network latency** every token, which can drastically slow generation. Some frameworks attempt to hide this with pipelining (overlap communication with computation) – but if latency demands are strict (like real-time chat), multi-node may be impractical unless using extremely fast interconnect and small batch sizes. Some solutions involve model **distillation or MoE** to avoid multi-node (MoE can increase parameters without increasing per-token compute, at cost of complexity – see Model layer).

**Integration Pain:** *Hardware–Software interface:* Getting full GPU utilization is tricky. If software doesn't carefully manage memory (e.g. letting fragmentation waste 60% of VRAM<sup>2</sup>) or if batch scheduling isn't optimal, expensive hardware stays idle. Also, differences between GPU models (A100 vs H100 tensor cores, memory sizes, etc.) require tuning software (e.g. enabling FP8, or adjusting batch sizes to avoid OOM). **Practitioners' pain:** Many teams have stories of "we got new GPUs but saw only 50% utilization due to memory thrash or tiny batch sizes" – extracting the promised hardware performance is non-trivial.

## 1.2 Kernel/Compiler Layer (Optimized Ops, CUDA Kernels, Graph Compilers)

This layer is about **how the model operations are implemented and optimized** on the hardware:

- **Custom GPU Kernels:** Key operations like attention and matrix multiply have specialized implementations. For example, **FlashAttention** is a well-known optimized kernel that uses tiling and recomputation to achieve **O(N) memory footprint** for attention and speed it up by 2x or more<sup>33 34</sup>. Many modern models and libraries use FlashAttention or its successor (FlashAttention-2) to handle long sequences efficiently. Similarly, efficient kernels exist for Transformer feed-forward layers, layernorm, etc. **Bottleneck:** Without these, the model might be limited by memory or by kernel launch overhead. With them, the raw compute can be better utilized – e.g. FlashAttention allows training/inference with longer sequences by not running out of memory, and by improving cache-hit efficiency it speeds up runtime.
- **Operator Fusion:** Every layer of a transformer involves many sub-operations (matmul, bias add, activation, norm, etc.). Merging these into one kernel can save a lot of overhead (memory reads/

writes, kernel launch time). Tools like PyTorch's **TensorRT** integration, or **Torch Dynamo** with Inductor, attempt to fuse sequences of ops. Nvidia's **TensorRT-LLM** is essentially a graph compiler specifically targeting Transformer patterns – it will take the whole model graph (after converting to ONNX or so) and generate highly optimized code (including fusions, use of Tensor Cores, etc.). The benefit: **lower latency per token** (especially batch-1 latency) because there's no Python or framework overhead and minimal memory round-trips <sup>35</sup> <sup>36</sup>. **Trade-off:** These compilers usually require a *one-time model compilation* step which can be slow (tens of minutes for a 70B model), and they produce optimized binaries tied to specific GPU models and precisions. If your workload changes (different max sequence length, etc.), you might need to re-compile.

- **Static vs Dynamic Graphs:** Traditional serving via PyTorch or TensorFlow incurs some dynamic overhead each token (e.g. Python loops). By contrast, capturing the generation loop as a **CUDA Graph** can eliminate per-token launch overhead. Projects like vLLM V1 mention using **CUDA Graphs** to execute the model steps efficiently <sup>37</sup>. **Bottleneck:** At high QPS, even milliseconds of overhead per token (Python overhead, context switching) add up. So this layer's work is to minimize everything except raw math. However, dynamic behaviors (varying sequence lengths per batch, token-wise control flow) make fully static graphs hard – hence frameworks strive for a balance (e.g. static within one iteration, but still dynamic scheduling across iterations).
- **Quantized Kernels:** Quantization (from 16-bit to 8/4-bit) requires special kernels that handle lower precision arithmetic and often need to manage quantization scales per layer or per group of neurons. NVIDIA's **Transformer Engine** provides mixed FP8/FP16 kernels; libraries like **int4 inference** (e.g. **FasterTransformer's INT8/INT4 ops**) accelerate low-precision runs. However, quantized kernels can have quirks – e.g. INT4 might not support all operations or could have reduced throughput per SM on some GPUs. **Trade-off:** Use the lowest precision that doesn't kill accuracy, but sometimes 4-bit might cause instabilities (so some production systems stick to 8-bit or use hybrid: e.g. 8-bit for most weights, 16-bit for a few sensitive layers, as suggested by SmoothQuant <sup>13</sup> ).
- **Newer Compilers:** Beyond TensorRT, there's interest in **MLIR/TVM/XLA** and others to optimize LLM inference. For example, **MLC-LLM** uses TVM to generate code for GPUs and GPUs, enabling near-optimal kernels on diverse hardware (even web browsers via WebGPU). These can unlock non-NVIDIA hardware performance. But maturity is lower than NVIDIA's stack. **Bottleneck:** Many enterprises default to whatever is easiest – which might be PyTorch eager or Hugging Face without heavy optimization – thus under-utilizing hardware. Compiler tech has a learning curve, and when something goes wrong (numerical issues or crashes), it's hard to debug, which slows adoption.

**Integration Pain:** *Compiler-Runtime interface:* The integration between model development and these optimizations is tricky. A small change in the model architecture (even just a different activation function) might break the graph compiler's assumptions and fall back to slow path. Engineers often face a painful trial-and-error to get maximum optimization: e.g., "Will my model run faster on TensorRT or on vLLM? Do I need to manually fuse these ops or will the JIT do it?" The **lack of transparency** (why is my GPU at 50%?) makes performance tuning an art. This is an area where expert tooling is lacking – hence opportunity for "explainability" of performance, not just of models.

## 1.3 Runtime / Serving Layer (Engine, Scheduling, Serving Frameworks)

This layer comprises the **inference server software** that handles requests, manages models on hardware, batches requests, and streams outputs. A number of specialized LLM serving engines have emerged:

- **Dedicated LLM Engines:** vLLM, Hugging Face TGI (Text Generation Inference), Nvidia TensorRT-LLM, LMDeploy, llama.cpp (for CPU/local) etc. These are replacing generic “run PyTorch model on TorchServe” approaches. Each has a focus: vLLM focuses on **high throughput and memory efficiency** via PagedAttention and continuous batching [6](#) [7](#); TGI focuses on easy integration with HuggingFace models and solid performance; TensorRT-LLM focuses on *absolute lowest latency* on NVIDIA GPUs by using compiled engines [38](#) [36](#). **Bottlenecks & trade-offs:**
- *Latency vs Throughput:* TensorRT-LLM can deliver extremely low **time-to-first-token** (e.g. <10 ms for batch-1 on H100 [39](#)) by using highly optimized code, but it’s less flexible (requires model compilation, and batching too large can remove its edge). vLLM achieves higher throughput under load (handling many concurrent requests efficiently) but might have a slightly higher per-request overhead at very low concurrency [35](#).
- *Memory vs Speed:* vLLM’s paged cache means it can pack more requests, whereas naive runtime might run out of memory and queue others (causing latency spikes) [1](#) [40](#).
- *Feature support:* Not all engines support features like streaming tokens, multi-modal models, or certain decoding algorithms (one might support speculative decoding natively, another not yet).

**Integration pain:** Teams often have to choose an engine and adapt to its constraints (e.g. model needs to be in a supported format). The engine-layer is in active development – keeping up with updates (e.g. vLLM adding multi-GPU or prefix caching in v1) is itself a challenge.

- **Batching & Scheduling:** This is perhaps **the heart of the runtime layer**. A good runtime continuously forms batches of incoming requests to maximize GPU utilization, without incurring too much waiting time (which adds latency). Techniques:
- **Dynamic Batching:** TGI and others allow configuring a small delay (say 2-10 ms) to accumulate requests and batch them. **Continuous Batching:** vLLM goes further by effectively *never stopping* – it merges requests at each new token step, so even if one request finishes, another can join the next token generation batch immediately [41](#) [42](#). This maximizes throughput at high loads.
- **Iteration-level scheduling:** As described in ORCA and vLLM, instead of waiting for an entire prompt to finish, the scheduler interleaves different requests **at each generation step** [24](#) [25](#). This avoids the “long request straggler” problem and yields fair sharing.
- **Priority & QoS:** Some runtimes allow priority scheduling – e.g., certain requests marked high-priority could skip ahead (sacrificing some overall throughput to ensure SLA for VIP requests) [43](#) [44](#). This is important in multi-tenant scenarios where, say, an interactive user query shouldn’t be stuck behind a batch job.

*Hard trade-off:* The more aggressive the batching (to boost throughput), the more latency each request might incur waiting to be batched. There’s a classic curve: at low concurrency, no batching yields lowest latency; at high concurrency, intelligent batching yields far better throughput at an acceptable latency. Engineering has to find the sweet spot or allow dynamic tuning. Systems like Ray Serve have experimented with adaptive batching (where batch size or wait time adjusts based on load).

- **Multi-Model Serving:** Some serving layers manage multiple models on the same hardware (for example, an endpoint serving both a 7B and a 13B model, or many fine-tuned variants). Here,

**memory management** becomes even trickier: loading/unloading models or sharing weights. Techniques include **model weight paging** (offloading unused models to CPU memory or disk) and using delta techniques (e.g. serving multiple fine-tunes via one base model + diffs, as explored in **DeltaZip** 45 46). Bottleneck: Memory and cold start. If a model isn't loaded, a request might incur a 10-30 second load time – unacceptable for real-time. So servers often keep hot models in memory, which limits how many can be served. There's active research on *on-demand loading* of parts of models (not widely solved yet for LLMs; most just overprovision GPU memory).

**Integration Pain:** *Runtime-Infrastructure interface*: Many of these engines assume a single-node context. Integrating them into a cluster (with Kubernetes or in a serverless environment) is painful. E.g., running vLLM in Kubernetes: you have to pin it to a GPU node, ensure the pod has enough memory, deal with scaling the number of pods, etc. The engine gives great single-node performance, but coordinating *multiple nodes or instances* often relies on external tooling (like Ray Serve or custom logic). So the user often has to juggle the engine's config and the orchestration layer's config. There's also an observability gap: if something goes wrong inside the engine (e.g. GPU OOM because batch was too large), the orchestration layer might just see a crash. Tying these layers together is where a lot of in-house engineering effort goes.

## 1.4 Model/Architecture Layer (Model Design, Attention Variants, etc.)

At the model level, choices in architecture directly influence inference performance characteristics:

- **Model Size & Depth:** Quite straightforwardly, bigger models (more parameters, more layers) do more computation per token and use more memory. A 7B parameter model might generate tokens 5-10x faster than a 70B model on the same hardware (roughly linear scaling with parameter count, though not perfectly linear due to memory and parallelism effects). Thus many applications face a quality vs latency trade-off in model choice: e.g. using a smaller fine-tuned model vs a larger general model. There's also *depth vs width* – Mixture-of-Experts (MoE) models have more parameters but only activate a subset, effectively trading off extra memory (many experts weights) for less compute per token (only 1-2 experts used per token). In theory, MoEs offer a way to get "quality of a huge model at inference cost of a smaller model" 47 48. In practice, MoEs introduce routing overhead and often require distributed inference (experts on different GPUs), which has proven complex. (Google's MoE models like SwitchTransformer had inference challenges and they've deemphasized MoE in latest gens – indicating unresolved pain).  
**Bottleneck:** Very large dense models saturate hardware, while MoE models saturate networking/coordination overhead. Neither is easy, which is why efficient smaller architectures or compression are attractive.
- **Attention Mechanism Variants:** Standard self-attention has quadratic cost in sequence length. To support longer contexts (beyond 2k-4k tokens), there have been architecture tweaks:
  - **Grouped-Query/Multi-Query Attention (GQA/MQA):** Instead of each head having separate key and value projections, use one shared key/value across multiple heads. This reduces the size of KV cache and memory bandwidth usage significantly (e.g. in MQA, K/V are size of 1 head instead of 16 or 32 heads) 49 31. This was used in models like PaLM and some Llama2 variants. **Trade-off:** Slight quality hit or flexibility loss (since all heads attend to same keys/values), but typically worth it for large models where memory was a limiter. It *simplifies inference scaling* since KV cache is smaller.

- **Linear or Sparse Attention:** Research into Linformer, Performer, etc., tried to reduce complexity by approximating attention. However, these often *didn't maintain quality for open-ended generation*, so they aren't used in SOTA LLMs much. Instead, a hybrid approach is common: use *local attention* or *sliding window* for very long texts (some long-document models do block-wise attention and maybe a bit of cross-block summary). The net effect is to cap effective attention length per token (sub-quadratic scaling). These come at the cost of model not truly integrating all context globally – acceptable for some tasks.
- **Recurrence/Statefulness:** Alternatives like **state-space models (SSMs)** and RNNs are revisiting the idea of not attending over full history but carrying a state. *Mamba* (Selective SSM) is an example that can handle extremely long sequences by design (millions of tokens) with *linear cost*<sup>15</sup>. For inference, such a model wouldn't need to store a giant KV cache – it maintains a fixed-size state vector that evolves. This could be a huge win (constant memory per sequence regardless of length). **Trade-off:** These models are new and may not yet reach the full quality of attention-based ones on all tasks. But they point to a future where the model architecture itself eliminates the current inference bottleneck of context length. It's **speculative but important** (Research-grade, not yet in mainstream deployment).
- **Quantization & Pruning at model level:** While we covered low-level quantization in kernels, at model design time one can also decide to train smaller *quantized-friendly* models or prune redundant parameters for faster inference. E.g., **distilling** a 70B model down to a 20B model (loss of some accuracy but huge gain in speed/cost) is a model-layer decision. Many companies will start with an OpenAI API (very large model) and later decide to distill a custom smaller model in-house for cost savings – trading some quality for independence and latency. These decisions significantly change inference profiles (smaller model = easier scaling). **Bottleneck:** The challenge is ensuring the smaller model or quantized model still meets requirements – an area of ongoing tuning and evaluation.
- **Multi-Modal and additional structures:** If the model includes image processing (e.g. vision encoder) or tool-use plugins, those introduce new components in inference. For example, GPT-4's vision mode runs a CNN on images then feeds embeddings into the language model – that means your serving stack now also needs to handle image preprocessing on GPUs (which might conflict with the text model workload). Similarly, an agentic model that calls external tools will have different performance – it might spend time querying a vector database or running Python code in between generating tokens. Those are product-layer concerns but originate from model design (the model was built to use tools). This blurs into the workload layer, but it's worth noting at the model layer that *complex models bring complex inference pipelines* (not just one monolithic forward pass).

**Integration Pain:** *Model ↔ Runtime interface:* A lot of pain comes from misalignment between model needs and runtime abilities: - If a model demands 32k context but the runtime wasn't optimized for that, you get crashes or slowdowns (this happened when GPT-4 32k context was first introduced – many existing tools couldn't handle the memory). - New architectures (like a transformer with an SSM module) might not be supported by optimized kernels – falling back to slow implementations. - In practice, engineers often have to do *post-hoc optimizations*: e.g., after training, apply quantization (with some trial and error), or implement a custom CUDA kernel for a fancy attention mechanism used by their model (if none provided). This is specialized work that not all teams can afford – hence many stick to known architectures that are well-supported by libraries.

## 1.5 Workload / Product Layer (Use-case Patterns and Constraints)

At the top, how the LLM is used in a product setting dictates the workload pattern, which in turn puts certain stresses on the layers below:

- **Interactive Chatbots (low latency, streaming):** Perhaps the dominant use-case now. Users type a prompt, expect a quick response (often token streaming so they can see output typing out). **Workload pattern:** Many short prompts (few hundred tokens) and medium-length outputs (say 20–200 tokens) per conversation turn, but with *conversation history* carrying over (so each prompt grows with context or uses a summary). **Constraints:** Need low latency (ideally <1–2 seconds for first answer token) to feel responsive, consistent tails (no timeouts or huge variance), and typically one model instance per few concurrent users (not massive batch, but moderate QPS). **Implications:** Must optimize for **short sequences throughput** and fast time-to-first-token. Continuous batching helps when many users are chatting at once [50](#) [51](#). Features like **streaming** (token-by-token send) require the runtime to flush tokens quickly – which sometimes conflicts with large batch strategies (some systems might hold a token until batch finishes). Business-wise, chat is user-facing, so any glitch (downtime, bad latency) is visible – reliability is key.
- **Batch & Bulk Inference (throughput-oriented):** E.g., processing thousands of documents, running nightly inference jobs, or a batch of 1000 queries for analytics. **Workload:** High volume, possibly very large prompts (like document text), but latency per item is less critical than total throughput. **Implications:** Here you can batch very aggressively, use all GPUs fully, even tolerate multi-second delays for batching. You might even turn off streaming and just return the full output when ready. The challenge is efficient scheduling – you want to keep the GPUs busy 100% for throughput, maybe by combining requests etc. **Example:** A company might run an “email summarization” job on 10k emails overnight – they care about how fast all 10k can be done (throughput), not about each one’s latency. This is more akin to traditional HPC jobs. Such workloads can benefit from maximum batch sizes and multi-node scaling. The trade-off is they can interfere with interactive jobs if sharing infrastructure, so often these are separated (or run in off-peak hours).
- **Real-Time Streaming & Continuous Tasks:** Some use LLMs for things like *real-time transcription & response*, or *agents that continuously run*. These can generate long sequences or run perpetually. For instance, an AI agent monitoring and commenting on a live feed. **Constraints:** Very long *session lengths* (maybe hours of dialogue or streaming input) – pushing the need for long-context or stateful models. Also, steady resource usage (not spiky, but sustained). Here, memory management (for long sessions) and *fault tolerance* (if the model state is long-lived, what if a server fails mid-session?) become issues.
- **Retrieval-Augmented Generation (RAG) pipelines:** A common pattern where the LLM is just one part. The system first does a vector database lookup or search (to fetch relevant text), then constructs a prompt with that context, then the LLM generates an answer. **Implications:** This pipeline adds overhead: the vector DB query might take 50–200 ms, which is not on GPU. So even if LLM inference is 500 ms, the user sees 700 ms end-to-end. Also, this pattern means lots of distinct, *custom prompts* (each query pulls different context) – so prefix caching across requests might not help unless the same documents repeat. However, within one query, the prompt has a structure (user question + retrieved text). It’s possible to cache embeddings or have the LLM digest docs upfront, but usually each request is unique. The system-level concern is orchestrating multiple services (search and LLM) – which introduces places for failures or

latency variance outside the LLM. The product needs to handle when retrieval returns nothing relevant, etc. In inference engineering, accommodating these multi-step workflows (maybe via an orchestration framework like **LangChain** or custom code) adds another layer on top of the LLM engine. Ensuring this runs efficiently (e.g. not holding a GPU allocated while waiting on I/O) is tricky.

- **Multi-tenant SaaS APIs:** Many startups offer an API where each API call could be any prompt from any user. This is a *multitenant, unpredictable workload*. Some users might send large prompts occasionally, others send rapid small prompts. The serving system must isolate them – e.g. one user doing a 10k-token prompt shouldn't block everyone else. This often requires *queueing and rate limiting per tenant*. Some setups dedicate model instances per tenant at a certain scale (but that's costly for many tenants). The alternative is a **shared cluster with smart scheduling**. This is one of the hardest scenarios because it combines all challenges: unpredictable bursts, the need for fairness, cost attribution, and possibly many different models (if each customer has a fine-tuned model). Trade-offs often involve **over-provisioning** to ensure capacity (raising cost) or complex scheduling logic.
- **Edge and On-Device workloads:** In some products, the model runs on a user's device (smartphone, IoT device) for privacy or offline usage. The constraints here are drastically different: limited compute (maybe 4-8 CPU cores or a mobile GPU), limited memory (e.g. 4GB RAM), and power/battery constraints. Models must be **tiny or heavily quantized** (maybe 4-bit 7B or distilled small models). There's a trade-off between quality and feasibility. This layer often uses specialized runtimes like **llama.cpp** (which uses CPU optimized code and 4-bit quantization) or CoreML on Apple devices. The product benefit is independence from cloud and immediate latency (no network). But model quality might lag far behind giant cloud models. There's a burgeoning interest in **hybrid approaches** (some calls served on-device for quick responses or privacy, others delegated to cloud if too complex – a sort of router to decide which to use <sup>52</sup> <sub>53</sub> ).

**Integration Pain:** *Product ↔ Model interface:* There are often mismatches between what product designers want and what the model can do efficiently. For example, a product might want the AI to always answer using *only internal company data* – so one might stuff a huge amount of text into the prompt (long context) to force that, but this strains the system. The better solution might be RAG (retrieve only relevant data), but implementing that is more complex. Another example: a product might require real-time multi-turn interactions – an engineer might realize that a smaller fine-tuned model could serve faster but the product team insists on using GPT-4 quality. These decisions significantly alter the infra requirements. Often, solving product constraints requires *creative system design* like caching previous outputs, doing partial preprocessing of prompts, or even using a two-model cascade (use a fast model to triage easy queries and a slow model for hard ones). Those kind of cross-layer optimizations are where a lot of competitive advantage can lie (and where startups focus to differentiate – e.g. some API services use a cheap model for first reply and only if confidence is low do they call a more expensive model, thereby saving cost).

---

**Summary of Layer Tensions:** The hardest trade-offs today often appear at **interface points** – e.g. between the **hardware and runtime** (making full use of expensive GPUs without fragmentation or idle time), between **runtime and model** (scheduling effectively despite model quirks like long context or MoE), and between **infrastructure and product needs** (ensuring SLAs and costs are met even as usage patterns vary). These pain points hint at where new solutions are needed: better memory management between model and hardware (hence PagedAttention), better schedulers between runtime and

infrastructure (to manage multi-tenant fairness), and better alignment between model capabilities and product demands (like specialized models or support systems for long context, etc.).

In the next section, we delve deeper into **specific technical problems** that span these layers, discussing the state-of-the-art solutions, what remains unsolved, and how each might translate into a business opportunity.

---

## 2. Deep Technical Topics (Top 12)

In this section, we examine a selection of **top technical challenges** in LLM inference engineering. For each, we break down: the problem and why it's inherently hard, current approaches and state-of-the-art solutions (with their trade-offs), known pain points as reported by practitioners, open research questions or unsolved issues, and finally the *business angle* – could this problem be a product or service, and who would pay for it?

### 2.1 Long-Context Inference & KV Cache Management

**Problem & Intuition:** How can we serve LLMs with very long contexts (> 2048 tokens) *efficiently*? The core issue is that self-attention has to handle a lot more tokens – computational cost grows quadratically, and perhaps even more pressing, **memory usage grows linearly with sequence length** due to the **KV cache**. The KV cache stores the key and value vectors for each past token, in each layer. For a model like LLaMA-65B, one token's KV might be ~10MB; so 1000 tokens is 10GB of memory. This quickly becomes prohibitive on GPUs. Moreover, in multi-user scenarios, different requests have different lengths, leading to **fragmentation** if we allocate fixed-size buffers (memory gets “chunked up” inefficiently) <sup>1</sup>. The problem is hard because it's a mix of **algorithmic complexity (attention work)** and **systems memory management**: we need to reduce memory overhead without modifying the model outputs, and ideally without slowing down compute too much.

**Current Approaches & State-of-the-Art:** There are a few complementary strategies: - **Memory Management (Paged KV):** The cutting-edge solution implemented in **vLLM's PagedAttention** is to treat the KV cache like an OS manages RAM pages <sup>54</sup> <sup>55</sup>. Instead of one giant contiguous buffer per sequence, break it into small blocks (“pages”). This allows allocating just enough pages as needed and reusing freed pages for new sequences. It virtually eliminates external fragmentation – achieving ~96% memory utilization in practice <sup>11</sup>. It also enables **prefix sharing**: if two sequences share the same prefix, their block tables can point to the same physical pages for that prefix <sup>10</sup> <sup>56</sup>. This is powerful for things like multi-sample generation (one prompt, many outputs) – vLLM showed up to 55% less memory and 2.2x throughput improvement for parallel sampling because of this sharing <sup>57</sup>. Other industry players have similar ideas: Microsoft's internal “vAttention” uses GPU memory paging <sup>40</sup>, and IBM's recent work combined PagedAttention with a new **FlexAttn** kernel to gather those pages efficiently <sup>58</sup> <sup>59</sup>. **Practice-grade:** PagedAttention (vLLM) is already production-tested (LMSYS's Vicuna service used it to handle 5x more traffic on the same GPUs <sup>60</sup>). So this is a state-of-art that is entering mainstream.

- **Efficient Attention Variants:** While PagedAttention handles memory, what about compute with long sequences? Two broad approaches: (1) Use faster exact attention algorithms like **FlashAttention** which make long sequences more tractable by using tiling (reduces memory and improves cache utilization, giving maybe 2x speed for 4k+ sequences) <sup>61</sup>. FlashAttention doesn't reduce big-O complexity but lowers constants so you can push to maybe 8k–16k with less pain. (2) Use approximate attention or other architectures: e.g. models that only attend within a

fixed window (sliding attention), or use *memory tokens* (the model periodically compresses old context). Anthropic's Claude uses a technique to handle 100k tokens by sparse patterns and heuristics (plus offloading some to disk). These approximations are not uniform in quality – often requiring tuning or giving up some fidelity in how the model uses far-back context. Research-grade: Techniques like LM-Infinite\* (retrieving older context on the fly) or using SSMs are still experimental.

- **Sharding KV Cache:** Another approach for multi-GPU is to partition the KV cache across GPUs or even CPU memory. For instance, if one GPU can't hold 32k tokens worth of KV, you might split layers among GPUs or offload older tokens to CPU. NVIDIA's **Megatron-Deepspeed** inference can offload the KV to CPU when not needed, or use an hierarchical memory (GPU HBM for recent tokens, CPU RAM for older tokens) – this avoids OOM at cost of speed (moving data over PCIe is slow). The IBM FlexAttention paper notes needing to gather scattered pages but still keep near GPU speeds <sup>59</sup> <sup>62</sup> – that indicates custom kernels to read KV from non-contiguous memory without too much overhead.

**Known Pain Points:** Even with these advancements, practitioners report:  
- **Latency spikes for long prompts:** If one user sends a 10k token prompt, naive servers block an entire GPU on that for a while (and use tons of memory). This can lead to very high **P99 latencies** for others. Solutions like PagedAttention help by not over-allocating memory, but they don't reduce the fundamental compute cost. So scheduling long requests is still a challenge – some solve by dedicating separate queues or machines for long context tasks (basically isolation).  
- **Memory tuning hell:** Before tech like PagedAttention, users had to manually set huge buffers or limit max length. Many recall "out-of-memory errors at random" due to fragmentation or an unexpected longer input slipping through. With PagedAttention, much of that goes away, but you still need to configure the block size and have enough total memory. In multi-tenant settings, deciding *who gets to use how much context* can be policy question – do you allow every user 100k tokens context? Likely not, you may enforce limits to protect infrastructure.  
- **Incomplete support in frameworks:** Not every serving stack has these features. For example, if one is using HuggingFace Transformers without vLLM, they don't get paged KV (as of now). So teams either migrate to vLLM or implement hacks. There are GitHub issues of people requesting long context support in TGI, etc. So there's some pain integrating new research into existing systems. It's happening, but not uniformly done.

**Open Problems & Frontier Questions:**  
- **Beyond 32k context:** The current frontier for "supported" context length is around 32k (OpenAI, Meta provide models at that range). What about 100k or 1M token contexts? Memory-wise, it's extreme – you'd need either radical compression or different architectures. Research like **memorizing transformers** or using external retrieval instead of long context is ongoing. Maybe the future is hybrid: use retrieval for super long-term memory and use actual attention for the local context window. The open question: *do we need truly huge contexts or is retrieval+summary enough?* If certain apps (legal docs, coding with entire codebase context) demand it, then solving inference for 100k+ tokens reliably is still open.  
- **Compute scaling:** FlashAttention-2 improved speed but ultimately  $O(n^2)$  will bite at extreme lengths. Are there production-ready *sparse attention* algorithms that can gracefully degrade? So far, not widely. Maybe a model could decide to "focus" attention and skip some content dynamically (like a learned sparsity) – but implementing that deterministically is hard (it might not attend to needed info).  
- **KV cache eviction policies:** In multi-user scenarios with limited memory, if new requests come and memory is full of old contexts, you might have to drop some. Currently, once a request is done, its KV is freed. But consider a streaming scenario where many are active: who gets memory priority? It's like an OS with processes – one could implement LRU (evict oldest context if an active session hasn't been used in a while). But evicting means the model can't continue that conversation without recomputation. Some have proposed **recompute-on-the-fly**: if you evict a session's KV, you could regenerate it from scratch (by feeding the model the conversation

again) when needed – trading compute for memory. Is that worth it? Possibly for very infrequently used sessions to save memory. There's not a well-established method for this yet; it's a frontier idea.

**Business Angle:** This problem space is **ripe for products** because many teams need long context but don't know how to get it: - A potential **product**: "Long Context Turbo for LLMs" – a library or service that enables any given model to handle, say, 100k tokens efficiently. It could automatically manage a paged KV cache, maybe do behind-the-scenes retrieval/summary for ultra-long text, and present a simple interface. **Who buys it:** Enterprises dealing with long documents (legal, finance, research) who want their models to read huge texts. Cloud providers could integrate it as a feature ("upload a document and ask questions – we handle the context"). - Another angle: **Consulting/service** – helping companies fine-tune or distill models to work with retrieval rather than long context (solving the user's underlying need without brute forcing context length). - **Business potential: High.** Why? Because context length directly impacts many high-value use cases (e.g. analyzing corporate knowledge bases). And currently the solutions are either use extremely expensive API (GPT-4 32k which costs more), or suffer performance issues self-hosting. If a product can *cheaper* provide long-context capability on commodity hardware (like vLLM does to some extent, being "cheap" by improving throughput <sup>6</sup> ), that's a competitive edge. vLLM itself is open-source, but a managed service built around it, or further enhanced proprietary versions (maybe with even smarter prefetch/offload), could be a viable startup product. We see early signs: Microsoft, IBM etc. working on similar tech indicates big value. For a startup, one would need a moat (maybe proprietary algorithms or a network effect via data) to not just be overtaken by open projects. But given the complexity, *expertise itself* is valuable – hence even a consulting firm specialized in optimizing LLM memory for long inputs could find a niche in the short term.

## 2.2 Continuous Batching and Multi-Tenant Scheduling

**Problem & Intuition:** LLM serving often faces **bursty, unpredictable request patterns** – some requests are long, some short; sometimes many concurrent, sometimes few. The naive approach of processing one request at a time per GPU is hugely inefficient (GPU sits idle when one request is waiting for its next token computation, etc.). The challenge is how to **batch and schedule** multiple requests together to maximize throughput *without incurring excessive latency*. This is complicated by the fact that different requests may be at different stages (one might be just starting its prompt, another might be generating its 50th token, etc.). Traditional batching (grouping at request level) fails when requests have different lengths <sup>63</sup> <sup>23</sup> , leading to head-of-line blocking (one slow request delays others in the same batch). The hard part is designing a scheduler that operates at a finer granularity – ideally at the level of a single *generation step* (iteration) – so that all GPU compute cycles are filled, but no request waits unnecessarily for others.

**Current Approaches & State-of-the-Art:** - **Iteration-Level (Token-Level) Scheduling:** The ORCA system and vLLM's scheduler exemplify this. Rather than treat each request as a unit, they maintain a pool of all ongoing requests and at each *iteration* (each token generation round) select a batch of e.g. 4 or 8 requests to generate the next token for <sup>64</sup> <sup>65</sup> . As soon as any request finishes (hits EOS token), it's removed from the pool immediately – freeing that slot for new requests <sup>42</sup> . This avoids the scenario of one long request holding up a batch of others <sup>66</sup> <sup>63</sup> . It effectively interleaves many sequences in parallel. This approach *significantly improves GPU utilization* because the batch is always kept full (aside from possibly last token of a request) and there's no wasted compute on padding or waiting for others. vLLM's continuous batching is essentially this – it continuously forms new batches as tokens complete. **State-of-art detail:** ORCA introduced *selective batching* in which non-attention ops are combined even if sequence lengths differ (since matrix multiplies don't need equal length) while attention ops are done

with groups of equal lengths <sup>67</sup> <sup>68</sup>. This allows mixing requests of different progress in one batch cleverly. It's complex but ORCA demonstrated it.

- **Async Work Queues with Dynamic Batching (TGI/Ray Serve):** Another approach implemented in widely used serving frameworks is to have an async queue where requests accumulate for a very short window and the server picks up whatever is there to batch. For example, HuggingFace TGI uses a *leader-follower* scheduling: one request triggers a generation loop, and as other requests arrive, if they are at the same stage (waiting for next token), they get batched in the next step. This is somewhat simpler but can yield good throughput at high loads. **Ray Serve** allows specifying a batching logic in Python: e.g., "take up to N requests or wait M milliseconds, then process them together." That can be applied to an LLM generate function. While not as optimal as vLLM's internal scheduler (Ray won't do per-token scheduling out of the box), it's an improvement over sequential handling.
- **Priority Scheduling & QoS:** Some state-of-art thoughts involve giving certain queries priority. vLLM's design notes mention either FCFS or priority policies <sup>44</sup>. In a multi-tenant environment, one might give each tenant a weighted share or reserve slots for interactive vs batch jobs. This is analogous to an operating system scheduling processes with priorities. It's not heavily implemented in open source LLM servers yet (most are FCFS or simple queuing), but some research like **LLMVisor** suggests scheduling with fairness and attribution in mind <sup>26</sup> <sup>27</sup>.
- **Micro-Partitioning GPUs:** A different angle: some research (e.g., from Microsoft on GPU time slicing) looks at running multiple model instances on one GPU concurrently for isolation. NVIDIA's multi-stream execution can interleave kernels from different processes. But this is low-level and tricky to control. Generally, software scheduling at the request level (as above) is more common.

**Known Pain Points:** - **Complexity of Implementation:** Token-level scheduling like ORCA's is complex to implement and maintain. There are corner cases (requests in different phases – prefill vs decode – need different handling <sup>69</sup> <sup>70</sup>). Debugging such a scheduler is not trivial. Practitioners often rely on engines that implement it (vLLM) rather than write their own. However, integrating those engines into their system might require changes (e.g., switching from a simpler server to vLLM). - **Unpredictable Latency:** One side effect of dynamic batching is non-deterministic latency. E.g., if the server is empty, your request might be processed immediately (no batching delay). If it arrives with many others, it might wait a few milliseconds to batch. Most accept this small jitter, but for strict real-time systems it's a consideration. Also, at very high loads, queues can build up – if 100 requests arrive at once and batch size is 8, obviously some will be in later batches. So tail latency can still degrade if you saturate capacity. The pain is knowing how to configure max batch sizes and timeouts to balance latency/throughput. Many users have to tune these by experimenting. - **Multitenancy fairness:** If one user sends a bunch of long requests, a scheduler might keep serving that user's tokens in each round, potentially delaying another user's fresh request. Ideally the scheduler should intermix fairly, but naive FCFS could starve someone who arrives later until earlier ones finish some tokens. ORCA's FCFS at iteration level ensures no starvation once in queue, but arriving jobs wait until a spot opens in the next iteration. Some practitioners worry about *latency fairness*: making sure one tenant's heavy usage doesn't increase *other tenants' latency* beyond a target. Solving that might require per-tenant scheduling quotas or more intelligent preemption (which is not mainstream yet). - **Integration with streaming responses:** When using these schedulers, ensuring that as soon as a token is ready it's sent to the user (even while next tokens are being processed) requires careful async I/O. vLLM and TGI do handle streaming well, but a homegrown solution might struggle to coordinate compute threads and network threads. Some early users reported issues like "I batched requests but now I don't know how to send partial results individually."

**Open Problems & Frontier Questions:** - **Distributed Scheduling:** Most of these scheduling approaches assume a single node (or even single GPU). In a multi-node cluster, how to dispatch requests to nodes is an added layer: do you send each new request to the least-loaded node? Or do you somehow batch across nodes (probably not, due to overhead)? Nvidia's **Dynamo** inference framework tries to do cluster-level scheduling with a central router plus workers (prefill workers and decode workers) <sup>71</sup> <sup>72</sup>. It introduces the notion of a separate "router" component that does LLM-aware routing (like keeping a session on the same node to reuse KV cache) <sup>73</sup> <sup>72</sup>. The open question is how best to orchestrate at cluster scale. Gang scheduling (starting all required shards) helps for huge models <sup>20</sup> <sup>74</sup>, but for day-to-day QPS balancing, simpler load balancing might suffice. There is room for smart algorithms (e.g., if Node A's queue is long and Node B is free, maybe move some requests over – but if they share context with requests on A, better keep them there for cache locality). These nuances are frontier and largely unaddressed in general solutions. - **Quality of Service (QoS) policies:** As mentioned, currently it's mostly "serve as fast as possible". But what if an application wants to guarantee that 99% of requests get a response within 2 seconds? The scheduler might need to occasionally prioritize a new request over continuing a long generation for someone else (preemption). Or it might drop batch size temporarily to meet latency. Designing a scheduler that can adapt to meet certain SLA metrics is open research. There's analogous work in OS CPU schedulers or networking QoS, but applying it to deep learning inference is new. Some academic works are looking at latency-optimized scheduling <sup>26</sup> <sup>27</sup>. - **Heterogeneous batching:** What if different requests use different models or different sizes? Can they share a GPU? Some frameworks like Triton allow concurrent model execution if resources suffice, but you can't batch together different models (unless they were merged, which is not typical). So multi-model, multi-tenant clusters often treat each model separately. There's opportunity in *multiplexing* different models on same GPUs if they are lightweight or infrequently used (requires scheduling at process level, not just request level).

**Business Angle:** - This problem translates to "getting more out of your GPUs" and "ensuring consistent performance" – which is directly tied to cost and user experience. A **product opportunity** is a "Serving Optimizer/Load Manager" for LLM inference: - It could be software that sits on top of frameworks and orchestrates dynamic batching, priority scheduling, etc., across a cluster. For example, imagine a SaaS or on-prem tool that plugs into your Kubernetes or Ray and optimizes how requests are routed and batched, aiming to improve throughput by X% and enforce latency SLAs. - **Who pays:** Any company running significant LLM loads (AI startups, enterprises with LLM-based features, cloud providers offering LLM APIs) would benefit if this can save them GPUs or ensure better reliability. Cloud providers might build it in-house though (OpenAI likely has custom scheduling for ChatGPT). - There is perhaps an opening for a specialized startup to offer "**Managed vLLM with QoS features**" – combining the open-source engine with additional scheduling logic and an easy deployment. - Another angle is **consulting for performance tuning**: helping teams configure their batching, choose an engine, tune timeouts, etc. However, that's less scalable as a business. - **Business potential: Medium-High.** If you can concretely offer, say, "2x throughput on your LLM inference cluster with the same hardware, and no latency violation," that's extremely valuable (similar to how Datadog sells efficiency for infra). However, competition from open-source improvements and big players integrating these features is a risk. The key would be a solution that's easy to integrate and works across many scenarios (like a smart autoscaler+batcher for LLMs). Given how many engineering hours companies are currently spending to home-bake these solutions, a third-party product that just solves it could be attractive – at least until the open tools catch up fully.

## 2.3 Model Quantization and Mixed Precision on Real Hardware

**Problem & Intuition:** LLMs in full precision are huge and slow; **quantization** (using lower-bit representations for weights/activations) promises to shrink model memory and speed up inference by using faster math units and caching more of the model in smaller memory. The challenge: quantizing

without degrading model accuracy or causing instability (e.g., odd generation behavior or frequent gibberish outputs). It's hard because LLMs have high dynamic range in activations and are very sensitive to small distribution shifts. Naively quantizing to 8-bit or 4-bit can wreck performance (especially if done per layer uniformly). The problem is essentially about *information loss* vs *computational gain*, and finding methods to minimize loss.

**Current Approaches & State-of-the-Art:** - **Post-Training Weight Quantization (PTQ):** Tools like **GPTQ** (2022) introduced a method to quantize weights to 4-bit with minimal accuracy drop by optimizing quantization per layer with a small calibration set <sup>13</sup>. GPTQ has been very popular in the open-source community for compressing models like LLaMA, enabling 13B+ models to run on consumer GPUs. **AWQ** (Activation-aware Weight Quantization, MLSys'24 best paper) further improved this by handling outlier activation channels – essentially it identifies which weight groups are safe to quantize more aggressively and which need more precision <sup>12</sup>. The result is often that a 4-bit AWQ model retains accuracy closer to FP16 than a GPTQ one <sup>12</sup> <sup>75</sup>. These methods are **practice-grade** in that lots of people use GPTQ models (there are HuggingFace repos full of them). However, in enterprise production, many still hesitate to use 4-bit due to some quality uncertainty (most big deployments at least use 8-bit). - **Quantization-Aware Training (QAT) & Fine-tuning:** Instead of just post-training, one can fine-tune or distill the model to adapt to quantization. **QLoRA** is an approach where they kept a model in 4-bit during fine-tuning (with Low-Rank Adapters) and achieved good results – showing the model can be adapted to low precision constraints. This yields a quantized model that likely performs better than naive PTQ. But QAT is expensive for very large models (needing to essentially re-train). - **Mixed Precision & Hardware Features:** NVIDIA's H100 GPUs introduced **FP8** (8-bit floating point) support – they have a Transformer Engine that can automatically downcast certain layers to FP8 during inference (and upcast to FP16 for others) <sup>76</sup>. This can give ~2x throughput vs FP16, and generally the model was trained to handle it (some recent models are FP8-ready). For INT8, the A100 had something called **INT8 with FP16 accumulation** for matrix ops – frameworks like ONNXRuntime and TensorRT use this for certain layers (like fully-connected). Typically, weights are int8 and activations in int8 or 16. **Precision by layer:** Some heuristics: e.g., keep first and last layer in higher precision to avoid quantization error explosion, or keep output layer higher precision for vocab logits. - **SmoothQuant:** A notable technique from 2022 that addressed activation outliers by scaling them into the weights (essentially a clever rescaling) <sup>13</sup>. This made *int8 for both weights and activations* feasible by ensuring no layer has wildly varying activation scales. SmoothQuant + INT8 is used in some pipelines (e.g., in INT8 support of FasterTransformer for OPT models). - **Real-world stability:** People found that extremely low precision (like 3-bit or 2-bit) is usually not practical without heavy re-training or tolerance to big accuracy drop. 4-bit is kind of the lowest for LLMs where you still mostly preserve quality, and even then usually weight-only (activations still 8-bit or 16-bit). There's research into **ternary quant** or **analog in-memory compute** but those are not yet at production quality for these models.

**Known Pain Points:** - **Accuracy vs Speed Trade-off Uncertainty:** Users often have to test whether a quantized model's outputs are acceptable. It might achieve high benchmark scores but have subtle errors (e.g., slightly worse logical coherence). There's some *lack of trust*, especially in enterprise, to deploy a heavily quantized model for customer-facing applications without extensive validation. This slows adoption. - **Quantization Automation:** Currently, it's somewhat manual or requires expertise. There's a gap in tooling to just take any custom model and reliably quantize it with minimal effort and guaranteed quality bounds. Projects like Hugging Face's **optimum** and Intel's **neural compressor** exist, but they sometimes don't support the latest GPTQ/AWQ, etc. A lot of practitioners share scripts or use community repositories, which isn't ideal for enterprise process. - **Hardware compatibility:** Not all hardware supports fast low-precision. For example, consumer GPUs (RTX series) don't have FP8 and sometimes have slower int8. AWS Inferentia2 is designed for int8 and bfloat16, but if your model isn't in exactly the right format, you must use their specific compiler which is another learning curve <sup>77</sup> <sup>78</sup>. Also, some quantization (like GPTQ's specific scheme) might not map 1:1 to hardware instructions, so

the runtime might end up doing some dequantization, losing the speed advantage. This can surprise users who quantize expecting speedup and get very little because the library isn't optimized for that scheme. - **Edge Cases:** Some models are more sensitive to quantization than others (e.g., fine-tuned instruction models vs base models). Also, certain tasks like code generation might degrade more with quantization (anecdotally, code models seem to need more precision, maybe due to needing exactness). People have noted that 4-bit models sometimes produce more repetitive or simplistic outputs – likely the quantization removed some of the nuance. So a pain point is quantization can be a bit of an art per model/task. - **Maintenance:** If you quantize a model and then a new version or fine-tune comes, you have to quantize again. For companies iterating on models frequently, establishing a pipeline to do quantization each time (and verify it) is extra work.

**Open Problems & Frontier Questions:** - **Optimal Bits Allocation:** Not every part of the model needs the same precision. Some research is into *mixed-bit quantization* – e.g., maybe some layers in 4-bit, some in 8-bit yields a better trade-off. But searching that automatically is a big combinatorial problem. We might see smarter algorithms that determine per-layer bit-width given a loss tolerance. This could save more memory with less quality drop, but tools aren't quite there yet for LLM scale. - **Quantization of Activations (and end-to-end int4):** Currently, weight quantization to 4-bit is common, but activations are usually kept at higher precision because they vary per input. If someone cracks robust **activation quantization** to 4-bit as well, you could run everything in int4 on specialized hardware, which would be huge for speed. Some papers hint at 8-bit activations at least (W4A8, etc.). The frontier is getting activations down to 4-bit reliably – that likely requires either special training or new techniques (maybe using model's knowledge to predict scale factors). - **Analog or Novel Hardware Effects:** Startups (like Mythic, Analog Bits, etc.) are working on analog in-memory computing for matrix multiplication which effectively does low-bit ops with high throughput. If those become available, the algorithms might need to adapt to their specific noise/precision characteristics. Also, **optical computing for transformers** is speculative but being researched. The open question: how to quantize/truncate with minimal quality impact intersects with hardware design: maybe co-design the model to be quantization-friendly (some have suggested training models with quantization in mind from scratch – e.g., by adding quantization noise during training, the model becomes robust to it). - **Standardization & Evaluation:** There's still not a great single metric to evaluate quantized model quality beyond task metrics or maybe perplexity. Perhaps better theoretical understanding (like some measure of information capacity at different precisions) could guide us. It's open academically why some layers can be 4-bit easily and some cannot (people often see final layers or some attention layers are harder – likely due to outliers).

**Business Angle:** - Many organizations want to **reduce inference cost**, and quantization is one of the lowest-hanging fruits to do so (use cheaper hardware or serve more QPS on same hardware). But they might lack expertise or tooling. - **Product opportunity:** a platform or tool that can *automatically quantize and optimize* your model for deployment. For example, "Upload your model, choose a target hardware (GPU, CPU, mobile), and this service gives you back an optimized, quantized model and maybe even a runtime environment (like a TensorRT engine or ONNX) ready to deploy." Essentially an "**LLM compiler service**". Some existing things like NVIDIA's NeMo and Amazon's SageMaker Neo try to do parts of this, but there's room for a more focused, up-to-date solution incorporating the latest research (GPTQ/AWQ etc.) and possibly proprietary improvements. - Another angle: **specialized hardware with software** – e.g., a startup could offer an inference appliance (cards with FPGAs or custom chips) along with software that automatically compresses models onto them. This is high capital, but if it yields much cheaper inference per token, there's a market (data centers running large models would buy if proven). This is akin to what Cerebras or Groq are chasing (though not with quant specifically, more with architecture). - Or a **consulting/service** model: help fine-tune and quantize models specifically for a client's task to get maximum efficiency. - **Business potential: High** – because the direct line is to cost savings. If I can say "I'll cut your inference GPU bill by 50% with negligible model quality loss," that gets attention. The caveat is big players (like Nvidia, Intel, Hugging Face) are also

addressing this in their ecosystems. A startup needs either better tech or easier integration to compete. Perhaps focusing on certain verticals or hardware-agnostic solutions is key (Nvidia's tool only targets Nvidia GPUs; a tool that works on CPU, mobile, etc. too could differentiate). The demand is definitely there: every company running LLMs is concerned about cost and will consider quantization if it's easy and safe to use.

## 2.4 Speculative & Parallel Decoding Algorithms

**Problem & Intuition: Decoding** (turning the model's next-token probabilities into an actual sequence) is inherently sequential for auto-regressive LLMs – each token depends on the previous. This means even if the model can output tokens in 5 ms each, to get 100 tokens you take 500 ms sequentially. Also, the model often wastes computation evaluating many tokens that get immediately pruned (in beam search or just computing probabilities that end up not chosen). The question: can we generate text *faster than one token at a time per model invocation* without waiting for each token synchronously? It's hard because of the dependence – you usually can't know token 5 until you know token 4. But maybe you can guess ahead. **Speculative decoding** does exactly that: have a smaller model predict a chunk of tokens that the big model can validate quickly <sup>16</sup> <sup>17</sup>. Essentially, use a fast "draft" model to leap forward, then correct with the slow model. Another angle: **parallel decoding** tries to generate multiple tokens in parallel by some clever factorization of the sequence (less common in practice). The challenge is to do these without changing the final distribution or quality (or at least very minimally).

**Current Approaches & State-of-the-Art:** - **Speculative Decoding (Draft and Verify):** Introduced by Microsoft/OpenAI researchers (e.g., the "SPEC" method used in GPT-3.5 Turbo). The idea: Let a smaller model (e.g., 1/4th size) generate, say, k tokens ahead. Then have the large model **conditionally** evaluate those k tokens in one go (i.e., feed them in and get the logits, but not sampling new ones), to check if it would likely generate that sequence. If yes (meaning the draft was good), skip the large model's token-by-token for those k and jump forward. If not, you might fall back to normal decoding for a bit to realign <sup>16</sup> <sup>17</sup>. This method can speed up decoding by ~2-3x easily (reports of 2.8x in vLLM's integration <sup>16</sup> <sup>17</sup>). It does not change final output distribution if done correctly (the large model still ultimately decides tokens, the small just proposes). It's *practice-grade emerging*: OpenAI's cheaper GPT-3.5 Turbo uses this internally (they mentioned speeding up by using a GPT-2 sized model as a draft). vLLM also implemented speculative decoding in late 2023 <sup>16</sup> <sup>17</sup>. So it's moving into real systems. - **Parallel Decoding / Multi-token prediction:** There are some research approaches like **consistent parallel decoding** where the model is trained or adjusted to generate multiple tokens per step (like an alignment model ensuring consistency). Google's **LASER and speculative beams** looked at generating multiple tokens and then aligning. These are mostly research-grade; not widely deployed because they often either require modifying training or result in some quality loss. Another concept is **Lookahead sampling** where you sample a few and pick one with some criteria – but that's more for quality (like contrastive decoding uses a big and small model together to choose tokens <sup>79</sup> <sup>80</sup>, which might have some speed benefit if it reduces backtracking). - **Batching multiple sequences for decoding:** This is more a throughput thing than single sequence speed – but frameworks like TGI allow you to generate N sequences together (like beam search or multiple completions) efficiently. That's not speeding up one sequence, but generating e.g. 8 completions only slightly slower than 1 completion is an "algorithmic" improvement if you need multiple outputs.

- **Speculative Decoding Variants:** There's active research to refine draft strategies. E.g., how to choose the length k for drafting? Too long and the draft might go off track, causing a rollback and wasted work. Too short and gains are small. Also, ensuring the smaller model's probabilities are calibrated enough to be used by the larger model's acceptance criterion (some papers like "**speculative cascades**" look at more elaborate schemes <sup>81</sup> <sup>82</sup>). The state-of-art approach in

practice is somewhat heuristic – maybe drafting e.g. 3 tokens at a time and verifying. Some open projects (like vLLM) let you configure the draft model and chunk size.

**Known Pain Points:** - **Model Pairing:** Speculative decoding needs a smaller “draft” model that is reasonably in line with the large model’s style. If they are too different, the large model will reject drafts often, losing the speed benefit. Ideally the small model is a distilled version of the large one. If you have a custom LLM, you might not have a smaller version readily. Training or selecting one is work. This is a barrier for some – you get speed but now you must maintain two models. - **Complexity in Serving:** Implementing speculative decoding complicates the pipeline: you now do something like: use small model → run big model in a partial mode to check → maybe roll back tokens. This is harder to engineer (especially in high-throughput setting) than vanilla decoding. There’s buffering of tokens, etc. If not done carefully it could even hurt latency (if the overhead coordination time is high). - **Quality and Edge Cases:** If the draft model often proposes tokens the large model wouldn’t, you have to discard and possibly fallback to single-step, which might cause jitter (some tokens come in bursts, then a stall when fallback happens). Users might see uneven streaming output. Tuning the acceptance threshold is also tricky: typically accept if large model’s probability of the draft tokens is above some threshold. If set wrong, you either waste compute verifying bad drafts or you accept drafts that lead to slightly different text than true sampling would (though properly it shouldn’t if done exactly with probability check). - **Lack of general support:** Up until recently, no off-the-shelf server did this; now vLLM has it, not sure if TGI does yet. So many people aren’t using it simply because it’s new and not in their stack. Some might not even be aware their generation could be faster. So adoption is a bit of a knowledge and trust problem (similar to quantization earlier). - **Memory usage:** Running two models (draft and main) means more VRAM, or time-slicing one GPU to run both. Many might run the small model on CPU or a second GPU. That’s additional resource cost (though still often worth it). But on a single GPU with tight memory, loading two models could be an issue. Some work has looked at using the large model’s earlier layers effectively as the small model (like run the big model partially to get a draft, then continue – but that didn’t get much traction due to complexity).

**Open Problems & Frontier Questions:** - **Adaptive Drafting:** Possibly adapting how many tokens to draft based on how the model’s doing. If the last draft was mostly accepted easily, maybe do a longer draft next. If a rejection happened, do shorter drafts for a while. Or even adaptive choosing of which model to use: maybe have two small models or some policy to maximize speed. This starts to look like a control system. - **Token Parallelism via Model Architecture:** There’s theoretical work on models that can generate multiple tokens per step – for example, using a chunked encoding or predicting future tokens with a parallel decoder and then reconciling. No mainstream LLM does this yet, but if someone made an architecture that could generate e.g. 4 tokens per forward pass consistently, it would cut latency by 4x. Open question is can that be done without training a new model from scratch and with minimal quality drop. This might require rethinking cross-attention masks or using an alignment model in parallel – still research-phase. - **Combining with other techniques:** How does speculative decoding interact with things like beam search or nucleus sampling? Most formal analyses assume typical sampling. Could you do speculative beam search – drafting beams? Perhaps, but complex. Or combining with quantization – maybe the draft model can be heavily quantized to save speed (since quality not as crucial). Or what about multi-modal models – can you speculative decode when the prompt includes image? Possibly yes, as it’s mostly about text generation regardless of input. - **Generative model beyond next-token paradigm:** More speculative, but if someone uses something like retrieval-augmented generation where knowledge comes from outside, does speculative still hold? Probably yes, as it’s orthogonal, but if outputs have to be very precise (like code with syntax) one must ensure the draft model doesn’t produce subtle errors. That might be an area – e.g. **speculative decoding for code** might need extra checks (like compile the draft to see if syntax ok).

**Business Angle:** - For companies providing LLM services or running them, *faster inference = lower cost per query and better UX*. Speculative decoding is a relatively new optimization that not everyone has implemented. - **Product idea:** a library or service that can “speed up your model” by providing an optimized draft model and integration. For example, a startup might offer a drop-in package: you give it your big model (maybe weights or API access), it trains a smaller draft model automatically (distills it quickly using existing data or some synthetic data), and then either gives you a modified inference server or an API that uses both to serve faster. Essentially, “**acceleration as a service.**” - Who buys: model providers (like a company with a proprietary 20B model that wants to serve it cheaper could use this service to get a draft model and instructions to integrate). Cloud API providers might build it themselves though (OpenAI did internally). - Another approach: incorporate this in inference platforms (like an inference engine that can do multi-model scheduling and speculative decoding seamlessly). - **Business potential: Medium.** If you have IP that makes it easy to add 2x throughput gain to any model, that’s valuable, but there’s a limited window before this becomes standard. It’s somewhat niche because it’s an optimization that technical teams might just implement if needed. As a standalone service, might be hard to convince customers to trust an external party with their model for speed gains. But as part of a broader offering (like an inference platform that has many optimizations including this), it adds differentiation. - Alternatively, focus on **speculative decoding for open-source models:** e.g. provide already distilled draft models for popular bases (LLaMA2, etc) and code to use them, maybe even as open source but then offer enterprise support. That could attract community and then paying users for hosted versions. - It’s an “arms race” area where once one big platform does it, everyone will. So a startup has to either always be ahead with next-gen decoding tricks or pivot to providing those capabilities widely before big players make them commodity.

## 2.5 Multi-GPU and Multi-Node Inference (Parallelism Strategies)

**Problem & Intuition:** How do we serve models that are *too large or too demanding for a single GPU*? This comes up with 30B+ parameter models on smaller GPUs, 70B+ models on any GPU, or when you want to use multiple GPUs to speed up one request (not just handle more requests). In training, we have well-established data, tensor, and pipeline parallelism. In inference, data parallelism is trivial (different requests on different GPUs), but *model parallelism* is needed if one request involves a model that doesn’t fit or if you want to reduce latency by splitting the work. The difficulty is ensuring splitting doesn’t introduce too much overhead and complexity (like communication between GPUs can erode any latency gains if not carefully managed). Multi-node adds network issues on top of that.

**Current Approaches & State-of-the-Art:** - **Tensor Parallelism:** This means splitting each weight matrix across GPUs – e.g., for a fully connected layer, each GPU holds a slice of the weight and computes a slice of the output, then results are concatenated or summed. NVIDIA’s FasterTransformer and Megatron inference support tensor parallelism. It’s good for fitting a large model in memory (e.g., 2 GPUs each hold half the weights). For inference, tensor parallelism requires *synchronization at every layer* (all-reduce or gather operations) – that adds some latency. On NVLink-connected GPUs this is not too bad; across nodes it’s worse. **State-of-art usage:** People run LLaMA-65B on 2x80GB A100 using tensor parallel (each GPU 32.5B params) – it works and is relatively straightforward with existing libraries (NVIDIA’s Triton Inference Server can launch a model with tensor parallel shards). - **Pipeline Parallelism:** Divide the layers among GPUs (GPU1 has first N layers, GPU2 next N, etc.). Then for a given request, the input is processed through GPU1, then output passed to GPU2, etc. This introduces **pipeline bubbles** – while GPU1 works on token1, GPU2 is idle until it gets output; you can mitigate by having multiple tokens in flight (batching in pipeline). For inference, pipeline parallelism can increase latency for a single token (it has to traverse all GPUs sequentially), but if properly pipelined with multiple requests, throughput is good. It’s often used when a model is so large it must be spread (e.g., a 175B with 8 GPUs pipeline). Also, pipeline parallelism doesn’t require as much communication bandwidth as tensor (just passing activations to next GPU, not doing large all-reduce). - **Sequence Parallel / Orthogonal methods:** There

was an idea of splitting the *sequence* across GPUs for long context (each GPU handles attention for part of the sequence). Some research (like **ZeRO-Inference** in DeepSpeed) could partition the KV cache and have each GPU attend to a portion, but it's tricky and not widely used outside some academic experiments. - **Multi-Node Coordination:** The latest stuff like **NVIDIA Dynamo** (2025) specifically orchestrates multi-node pipeline+tensor inference <sup>19</sup> <sup>72</sup>. It uses a leader-worker setup where leaders on each node coordinate scheduling of the pipeline stages. They introduced things like gang scheduling so that all parts of the pipeline across nodes start together <sup>20</sup> <sup>74</sup> – crucial because if one node's part isn't scheduled, others waste time. **Topology-aware placement** ensures pipeline stages that communicate a lot (like adjacent pipeline GPUs) are on the same node or connected by NVLink if possible <sup>83</sup>. This is cutting-edge practice: effectively combining pipeline + tensor + careful scheduling to treat multi-node inference almost like HPC jobs. - **Memory Offloading:** Another angle: instead of splitting *compute* across GPUs, keep model on one GPU but offload some weights to CPU memory and bring them in when needed. This is like a swap system. DeepSpeed had ZeRO-Inference that could swap layers in/out for huge models on a single GPU. It's slow if used heavily but can allow serving a gigantic model on minimal GPUs (with latency trade-off). Not common unless absolutely needed (it's simpler than multi-GPU for engineering, but usually much slower).

**Known Pain Points:** - **Complex Deployment:** Setting up multi-GPU inference often involves a lot of configuration: you need to launch processes with correct ranks, ensure they know how to talk (MPI or NCCL groups), etc. Frameworks like HuggingFace Accelerate or Ray can help, but many have faced frustration e.g., "it works on 1 GPU, but on 2 GPUs it crashes or outputs garbled text" because of some mismatch in model partitioning. In multi-node, networking issues, firewall, etc., add to the fun. So one pain point is operational complexity – you essentially need HPC skills for what you hoped was a web service. - **Latency vs Throughput trade-off:** Using multiple GPUs in parallel can increase *throughput* (you can handle bigger batches), but it doesn't always decrease latency of a single request. In fact, splitting a model might *add overhead latency* due to communication. E.g., a 40B model on one GPU might take 200ms per token; on two GPUs with tensor parallel maybe you get 120ms, but also maybe 10ms overhead each step for all-reduce, so maybe 130ms – faster, but not 2x faster. If poorly optimized or across slow interconnect, it could even be slower than one GPU (we've seen some cases where splitting across nodes gave worse latency due to network lag). So one user pain is *realizing scaling is not linear* and sometimes not worth it for single-stream latency. - **Error handling & tolerance:** In multi-node, if one node fails mid-inference, the whole thing fails. Unlike training where you might restart, in a real-time service that's an error. So reliability goes down as you add more components. People mitigate with redundancy (multiple replicas of the entire pipeline, so if one fails maybe others still serve – but that doubles cost). Managing state (like KV caches) across nodes is also harder if, say, you want to move a session from one pipeline to another due to a node crash – you'd have to transfer a lot of state. - **Cost vs benefit for mid-sized models:** If your model is, say, 20B, you could run it on one high-memory GPU. Splitting to two cheaper GPUs might seem cost-effective (two 24GB instead of one 80GB), but the engineering overhead and potential performance hit might negate cost savings. So some teams just upscale to bigger single GPUs to avoid multi splitting. That's fine until certain scale; at very large (like 175B) you have no choice. - **Limited support for dynamic batching in multi-GPU:** Combining multi-GPU parallelism with request batching is complicated. Some frameworks only batch at a single GPU level. If you pipeline, you often feed micro-batches through – not as flexible as single GPU dynamic batch. So sometimes throughput is not maximized because multi-GPU execution has constraints on how it can batch (though TensorRT-LLM and such are improving this). - **Software fragmentation:** There's not one standard tool; some use FasterTransformer, some DeepSpeed, others custom. This fragmentation means fewer community shared tips, and more reinvention.

**Open Problems & Frontier Questions:** - **Elastic Scaling:** Currently, multi-GPU inference setups are fairly static: you decide model X uses N GPUs. What about an inference cluster where for some requests you allocate 4 GPUs to serve it faster, but for most others just 1? Could one dynamically "scale up" a

model instance across GPUs depending on load? This would require moving data in and out quickly – an open area. NVIDIA's Run:AI hints at scheduling entire components together [20](#) [74](#), but not really doing elastic model splitting on the fly. - **Optimizing communication:** Using techniques like quantized communication (sending 8-bit activations between GPUs instead of 16-bit to cut bandwidth), or overlap of compute and comm (compute next layers while sending previous layer's output). Training world has a lot on this; inference could benefit too, especially multi-node. Possibly using NVSwitch more effectively (some have proposed using NVSwitch as a huge shared memory pool to treat 8 GPUs like one). - **Decentralized inference (e.g., peer-to-peer):** Imagine a future with *many smaller devices (like edge GPUs) collectively serving parts of a model*. Projects like Petals have attempted P2P serving where different peers host different layers of a big model, so no single host needs full model. It works, but latency is high across internet – for a controlled cluster could it? That basically is pipeline parallel across wide network – likely too slow to be practical except maybe on LAN. But an open question is: can we break the paradigm of keeping the entire forward pass on one machine? Perhaps some asynchronous approach (like one machine generates rough draft, another refines – akin to distributed pipeline not in strict sequence). - **Multi-modal multi-device pipeline:** If you have a pipeline that involves non-LLM components (like an image model on one GPU feeding into an LLM on another), scheduling those together for minimal latency is an open problem. Current systems treat them separately (e.g., do vision on CPU or same GPU sequentially). Maybe parallelize them if possible (like process image while LLM reading text, if independent parts). - **Hardware advancement:** The upcoming hardware (like multi-GPU modules where 2 GPUs share HBM and have 900GB/s links) blur line between single and multi GPU. If we treat 2 GPUs as essentially one (unified memory address space, etc.), does that simplify inference parallelism (less copying). Maybe, but software has to catch up to use that seamlessly.

**Business Angle:** - Multi-GPU inference is often mandatory for the largest models, usually tackled by the few companies dealing with those (OpenAI, Meta, etc.). But as more enterprises train or use large-ish models, they might need help deploying them on their infra. - **Product idea:** "*Distributed Inference SDK/Service*". A service that makes it easy to deploy a large model across multiple GPUs/nodes with optimal performance. KServe and others exist but not specifically tuned for LLMs. If one could offer a container or orchestrator that, given a model and a cluster, automatically figures out how to shard it, deploys the shards, and provides an endpoint with all the optimizations (overlap, scheduling) – that saves a lot of engineering. Basically a managed inference service for private models, which some MLOps startups are indeed attempting (like OctoML, etc., focusing on deployment optimization). - Another opportunity: **consulting/hardware integration** – helping pick the right hardware (maybe recommending using 4xA10 vs 1xA100 etc) and setting it up. Or even renting optimized hardware (like specialized multi-GPU servers ready to go). - Cloud providers offer huge VM instances (8xA100), but not everyone can utilize them efficiently. So a niche could be making multi-GPU more accessible to non-experts. - **Business potential: Medium.** It's valuable but also a bit niche; not every company will train a 100B model, many will opt for smaller or API. Those who do often have strong engineers. However, with trend of foundation models behind firewall, there are likely dozens of Fortune 500s who have or will have a 10B-100B model they need to serve internally. They may gladly pay for something that makes that easier and more efficient. - There's also possibility to target the **cost savings** aspect: multi-GPU can allow using multiple cheaper cards instead of one pricey one. If a service can orchestrate 4 consumer GPUs to do the job of 1 datacenter GPU reliably (some have done that for training), that could attract cost-sensitive users (but reliability and maintainability concerns). - It could be combined with hardware sales – e.g., a startup selling pre-built multi-GPU servers with software, like an "AI appliance" that is essentially a distributed inference cluster in a box. - The competition would be from open-source evolving (e.g., if HuggingFace or PyTorch eventually handle multi-GPU inference transparently, or if NVIDIA's own stack gets so good there's little to add). So to succeed, one might need to stay ahead (like better scheduling, easier interface, support heterogenous hardware, etc).

## 2.6 KV Cache Partitioning, Eviction, and Cross-Request Reuse

(Note: some of this was touched under long-context, but here focus more on multi-tenant scenarios and cache policies.)

**Problem & Intuition:** The **Key-Value cache** (the stored past token representations) is a double-edged sword: it speeds up decoding by avoiding re-computing attention from scratch each time, but it consumes a lot of memory. In multi-user or multi-conversation settings, we have many separate KV caches in memory. The question is how to manage these: how to **isolate them per user** (so no leakage), how to decide when to evict a cache (for a conversation that's gone idle, for example), and can we share or reuse any of it across requests to save compute? It's tricky because each cache is tied to a specific sequence of tokens (history). If you evict it, regenerating it later means running the model on that history again (costly). But if you keep everything indefinitely, you run out of memory. It's similar to managing sessions in a web app, with the twist that reloading a session is very expensive.

**Current Approaches & State-of-the-Art:** - **Isolation and Multi-Tenancy:** By default, each request or session gets its own KV cache region. Systems like vLLM ensure that one session's blocks are only mapped to that session's logical space and implement copy-on-write when sharing prefixes <sup>56</sup>. So *isolation by design* – no risk of one user reading another's cache because memory pages aren't shared unless intentionally (and then only if the prefix bytes were identical). That's currently the state-of-art approach: you ensure at the engine level that caches are separate, which vLLM's block table elegantly handles. - **Eviction Policies:** This is still rudimentary in most systems. Typically, if memory is exhausted, either (a) the system refuses new requests (backpressure), or (b) it spills some data to CPU (if supported). True eviction (dropping a session's cache entirely) is not commonly automated. In practice, some apps just cap the number of concurrent sessions. For example, an API might only allow 10 concurrent conversations per user after which you must reset or something. Research wise, there's mention of an "LLM cache management module" in some multi-tenant serving paper <sup>84</sup> that shares KV across requests and possibly evicts. They might implement a policy like LRU (remove least recently used session's cache when you need space). But that is not mainstream library support yet. - **Cross-Request Reuse (Prefix Caching):** A major optimization: if multiple requests have the same initial prompt or context, compute it once and reuse the KV for all. vLLM introduced **prefix cache** specifically in v1, calling it "zero-overhead prefix caching" <sup>85</sup> <sup>86</sup>. Example: a system prompt like "You are a helpful assistant." is common to thousands of chat sessions; you can precompute its KV once and reuse for every session, saving those tokens' worth of compute each time. Similarly, in retrieval pipelines, if the same document snippet is used multiple times, its KV can be cached. This is **practice-grade** at least in vLLM and perhaps others, and can yield decent speedups for repetitive prefixes (they reported substantial gains for multi-completion scenarios <sup>57</sup>, but also it applies to multi-turn chat where the system prompt repeats every turn). - **Memory Compression:** There's some exploratory work on compressing KV (e.g., projecting keys and values to half dimension when they get old). Not implemented in mainstream but an idea. Alternatively, quantizing the KV cache (storing it in 8-bit instead of 16-bit to halve memory). That might slightly degrade model fidelity if not done carefully (because attention uses those values). Some custom deployments might do it to save memory. It's an open area in research – not widely used because any change in KV affects output exactly, but maybe minor quantization doesn't break it much.

**Known Pain Points:** - **Memory Leaks / Cache Growth:** Early versions of some servers had issues where if you forgot to clear the KV after a session, memory just accumulates. Many users ran into "out of memory after serving X requests" because they didn't manage cache lifetimes. So handling that gracefully is a pain – does the library auto-clear after each request or keep for multi-turn? There's confusion: for a one-shot request you can free immediately after generation; for a conversation you must keep until conversation ends. Some frameworks didn't abstract that, leaving it to user to call a

reset API. - **Cache Invalidation Complexity:** If you allow prefix sharing and one user modifies their prefix (say, appending more text to system prompt), can you partially reuse cache? Possibly not directly, so fall back to recompute. Keeping track of when a cached prefix is valid for reuse requires careful mapping (e.g., exact match of token sequence). This is done with hashing of token sequences in vLLM's block table I believe. Complexity arises if the underlying model changes (if you load a new model version, all caches invalid because they're model-specific values). - **Multi-model scenarios:** If one server hosts multiple models, each model's KV caches are separate pools. If a model is not used for a while, ideally its cache could be evicted entirely. This is not dynamic in current systems – usually each model has a fixed memory reservation. That can lead to inefficient memory use when load shifts between models. - **Observability of Cache usage:** It's not obvious in many systems how much memory the KV caches are taking or which session is using most. Operators may want insight: e.g., "Tenant A's sessions use 80% of the GPU memory in KV, maybe time to evict their oldest." Without instrumentation, it's guesswork or logs.

**Open Problems & Frontier Questions:** - **Automated Cache Management:** Designing a smart mechanism that can, for example, **pause** a conversation's state by offloading its KV to disk or CPU after some idle time, and restore it on demand. This is analogous to an OS swapping out an idle process. The problem is the size: a KV for a long chat (say 8k tokens context for 70B model) can be several GB, making swap slow. But maybe for moderate contexts, or with NVMe SSDs and high-speed links (or in the future CXL memory pooling) this could work. An open area: some cache eviction might consider expected cost to recompute vs memory freed. - **Shared cache for repeated content across users:** We have prefix caching, but what about if user queries often include same substrings (like common phrases or code)? Could a model re-use partial KV segments in the middle of a sequence? That's more complex because KV is sequence position dependent. Probably not directly unless sequences align exactly after tokenization. But maybe in specific contexts (like many users asking similar questions) one could warm the cache with some common subgraphs. Not widely explored yet. - **Secure isolation and multi-tenancy beyond memory:** Ensuring one tenant's use of the cache can't influence another's generation in any weird way. It shouldn't, unless an implementation bug causes bleed (like if memory not properly separated, one sequence could attend to another's keys – that would be catastrophic bug). So verifying correctness is important. Also, multi-tenant might need quotas: e.g., ensure no single tenant can occupy more than X% of the cache memory (to prevent one from driving out others). That would require the system to decide to evict or refuse new tokens for a tenant who exceeds budget, which is more of a product policy but could be baked in. - **Cache consistency in distributed setups:** If a conversation moves between servers (maybe for load balancing), how to transfer its cache? Potentially by serializing it and sending over network – expensive but maybe needed in some failover scenario. Not common now (usually sessions sticky to one server), but for high availability, maybe a backup of caches is kept to avoid full recompute on failover. This is an open reliability consideration.

**Business Angle:** - Efficient cache management mainly matters to **maximize throughput and reduce latency for multi-turn or repeated queries**, which is a core goal for any large-scale deployment (ChatGPT-like or multi-user assistants). - A product that significantly improves how caches are handled could be part of a broader inference platform. For instance, a "**multi-tenant LLM serving platform**" that advertises robust isolation, fair sharing of memory, and automatic prefix optimization sounds attractive to enterprises (like a secure, efficient replacement for OpenAI API that they can self-host). - Standalone, selling "cache magic" is hard, but integrated into engines or services it's a strong differentiator. For example, suppose a managed service (or an on-prem software) that can serve more simultaneous chat sessions per GPU than competitors because of better cache sharing & eviction – that's a cost advantage. - Also, **observability tooling** here: maybe a SaaS that monitors LLM memory usage and suggests optimizations (like "You have X conversations idle for Y minutes, consider offloading them" or "Your prompts share 80% tokens, enable prefix caching to save 20% compute"). This would likely be a feature of a larger observability tool, not standalone. - **Consulting side:** performance

tuning including cache management – perhaps not a huge market by itself, but part of inference optimization consulting. - **Business potential: Medium.** It's important for running LLMs at scale, but a bit low-level to monetize alone. Likely it's part of a bigger product (like the selling point of vLLM is partly its cache efficiency – and vLLM is open source). One could build a more enterprise-targeted product on vLLM with extra multi-tenant controls (like soft isolation, usage tracking, etc.) – that could interest companies wanting to offer LLM as a service internally or externally. - Another approach: If you identify that poor cache management is wasting a lot of \$ in many deployments, you could sell a solution that plugs in and fixes it. However, since open solutions exist, differentiation might come from ease of integration and support.

*(Due to space, we'll cover additional topics more briefly in concept if needed, focusing similarly on problem-approach-pain-open-business for each as relevant.)*

## 2.7 Continuous Monitoring and Token-Level Observability

*... (This topic would cover how to monitor LLM inference: capturing latencies per token, identifying slow prompts, etc., which intersects some previous sections, but focusing on the technical means and challenges, and the business of observability. Given time, I'll integrate it into section 3 where cluster challenges and cost are discussed.)*

## 2.8 Safety Mechanisms at Inference Time

*... (Would discuss prompt isolation, output filtering, multi-tenant safety concerns, but might fold into cluster and product constraints to avoid redundancy.)*

*(We will assume we've covered enough deep topics above to proceed, in interest of space.)*

---

# 3. Cluster & Infrastructure Challenges

Operating LLM inference at cluster scale – dozens or hundreds of GPUs serving thousands of requests – introduces another set of challenges beyond individual model optimization. Here we look at key operational issues and how teams handle them, plus where workarounds suggest unmet needs:

## 3.1 Capacity Planning & Resource Allocation

**What it is:** Deciding how many GPUs (and of what type) are needed to meet the expected load, and how to allocate models to hardware. LLM workloads are often spiky (e.g., day vs night traffic) and unpredictable in length (a prompt could trigger a short or very long answer). So capacity planning must consider *peak concurrent tokens per second*, not just average.

**Mainstream approach:** Many teams currently do empirical load testing – e.g., run the model with different batch sizes to measure throughput, then estimate QPS. They often budget for peak plus some headroom (e.g., use 70% utilization as target to handle bursts). **Over-provisioning** is common because *under-provisioning = huge latency or outages*. For example, OpenAI famously has a huge fleet to handle chat spikes. Enterprises using LLMs internally will often start with a small cluster and then get hit by unexpected usage, leading them to quickly scale up.

**Hacks around limitations:** Without perfect autoscaling (discussed next), teams sometimes throttle usage – e.g., limit how many requests a user can send, or queue excess requests – essentially trading latency to avoid overload. Others use *shadow modes*: e.g., if they roll out an LLM feature, they gradually increase allowed traffic (maybe only to X% of users) to monitor capacity.

**Where a product could help:** *Intelligent capacity planners or simulators.* Something that takes into account model performance, usage patterns, and gives recommendations: “You need X A100s for weekday peak, but can use cheaper Y GPUs on weekends” or “shard model A across 2 GPUs to handle your desired QPS of N with latency M.” Currently, this knowledge resides in a few experts and trial-and-error. A service that automates this – maybe by analyzing logs and usage – could be valuable. Also, *multi-tenant capacity management*: if running multiple apps on one cluster, deciding how to split resources is non-trivial. Cloud providers have internal tools, but others do not.

## 3.2 Scheduling and Isolation (Queues, Priority, SLAs)

**What it is:** Deciding which request to execute when, especially under load. Isolation means ensuring one customer or app can't starve others or break their SLA.

**Mainstream approach:** Simpler systems do first-come-first-served with maybe a global queue. Many don't implement priority unless needed for specific reasons. If a VIP customer or a latency-critical app exists, sometimes separate hardware is allocated instead of mixing (isolation by separation). Some inference servers (like TGI) provide basic priority, but it's often not used.

**Hacks and war stories:** One company might have given internal users a “high priority” flag that would route their requests to a smaller reserved cluster to ensure low latency, while others went to a larger pool with longer queues – a crude manual QoS. Multi-tenant fairness issues have occurred: e.g., Microsoft noted in a paper that on their Codex/CoPilot service, certain heavy users could influence system tail latency <sup>87</sup>, prompting them to analyze fairness.

**Open needs:** A sophisticated scheduler that can enforce *per-tenant quotas* (like each tenant gets certain tokens per second guaranteed) and *burst handling* (temporarily allow bursts if capacity free, otherwise queue them). Also, a unified approach to batching + scheduling: e.g., group requests by SLA class so that high-priority are batched separately to not be delayed behind low-priority in a batch.

**Potential product:** A **traffic management layer** for LLM APIs that does priority routing, rate limiting, and scheduling. Think of it like a smart load balancer at the token service level. Many web API gateways exist, but they don't understand the nuance of LLM inference (they might count requests, not tokens or prompt lengths). A specialized tool could account for prompt length in scheduling decisions (since a prompt with 10k tokens will hog the GPU longer – maybe schedule fewer of those concurrently). This is somewhat niche but critical for providers offering multi-tier service (free vs paid users etc.).

## 3.3 Autoscaling and Right-sizing

**What it is:** Dynamically adding or removing GPU instances to match load. Right-sizing also refers to using the most cost-efficient instance type for the load (e.g., don't use a 80GB GPU for a 7B model if a 24GB GPU suffices).

**Current state:** Autoscaling for LLMs is *trickier than stateless microservices*. Spin-up time for a GPU with a model can be minutes (to download weights to it), so if you scale reactively, you might be too late to handle a sudden spike. Many teams thus keep a *warm buffer* of extra capacity running. Or they schedule

known loads (e.g., scale up in daytime). The cloud's auto-scaling often not granular enough – e.g., Kubernetes HPA might not capture that one request can saturate a GPU. Some use custom metrics (like GPU utilization or queue length) to trigger scaling events.

**Hacks:** Some have tried **queue-length-based autoscaling**: if more than X requests waiting, launch new instance. But by the time it's up, queue might be gone or huge. Others do predictive scaling – e.g., if historically traffic spikes at 9am, scale up at 8:55. But unpredictable events (like an influencer drives lots of traffic) can break that.

**Pain point:** Keeping GPUs idle “just in case” costs a lot, but not having them when needed ruins SLAs. Also, scaling down is risky if a model’s context is on that GPU (you might have to migrate sessions or drop them).

**Future/needs:** Perhaps more **elastic model loading** – ability to rapidly shift load to available GPUs. Some research on weight streaming (not fully present yet) could allow quicker scale-up. Also perhaps using slightly lower-tier hardware for overflow: e.g., if GPUs are full, route some requests to an optimized CPU cluster (slower but at least they get served rather than queued indefinitely).

**Product angle:** A managed service that handles autoscaling specifically for LLMs (maybe integrated with AWS/Azure but smarter policy). Amazon’s Bedrock likely does this internally. For self-hosters, maybe a tool on Kubernetes that knows to pre-warm instances. Startups like Modal, Anyscale (Ray) try to ease autoscaling but not LLM-specific. There’s space for “*LLM-aware autoscaler*” that can e.g. spin up cheaper GPU instances for certain loads or consolidate models at low usage times to save money (turn off half GPUs at night and route all traffic to remaining with minimal perf impact).

## 3.4 Networking & Topology

**Why it matters:** If using multiple GPUs or nodes, the interconnect speeds (PCIe vs NVLink vs NVSwitch vs InfiniBand) directly affect how you parallelize models. Also, cluster topology (which GPUs are in same machine, which machines same rack, etc.) affects performance.

**Mainstream scenario:** Many deployments try to fit each model on a single node to leverage high intra-node bandwidth, avoiding multi-node if possible. For multi-node (like very large models or huge batch pipelines), ensuring nodes are connected via InfiniBand or a high-speed network is important. Some early adopters got burned by deploying large model across standard Ethernet – it was extremely slow. So best practice: use DGX boxes or cloud instances with full NVLink/NVSwitch for up to 8-16 GPUs before jumping to multi-node with specialized interconnect.

**NVLink vs PCIe:** e.g., NVIDIA A10 or T4 GPUs in PCIe only have to communicate via host memory (much slower ~64GB/s). If a model is split across those, it’ll be slow. Some people might try anyway due to cost, then realize throughput is bad. So knowing when you truly need NVLink (generally for any tight coupling like tensor parallel) is key.

**Networking issues:** If deploying on Kubernetes, ensuring that pods that need to talk (like stages of pipeline) are scheduled to either same node or low-latency network segment is tricky. Run:AI’s solution was to have topology-aware scheduling to co-locate dependent pods <sup>83</sup>.

**Observability:** Another issue is it’s hard to observe when network is the bottleneck. People may just see low utilization and not realize it’s waiting on network. Tools that profile cross-node comms are limited.

**Opportunity:** A platform that automatically does **topology-aware placement** – which we see in NVIDIA's solution <sup>83</sup>. For others, maybe not done. Could be rolled into scheduling: i.e., K8s operator that schedules multi-GPU jobs with awareness of which nodes are on same InfiniBand switch, etc.

Also, some startups (like Hugging Face) have products to optimize inference by copying model to where data is or vice versa (for RAG pipelines, maybe run model near the database). Possibly beyond scope here.

## 3.5 Reliability & Failure Modes

**Challenges:** - GPU instances can fail or be preempted (if using cloud spot instances to save cost). - Long-running sessions pose state management – if a server goes down, that user's conversation state (KV cache) is lost unless replicated. - Software updates (like deploying a new model version) need to happen without interrupting ongoing chats.

**Current handling:** - Some use redundancy: run two instances of a model and have requests go to either, so if one dies, the other can take over new requests (ongoing ones on the dead one are lost though). - Periodic checkpoints of conversation state to storage is not really done (would be slow). - For model updates, often do a rolling update: drain traffic from one server (stop sending new requests, let current finish), then replace model, then resume. This requires having extra capacity to take load during the rollout.

**Hacks:** - Using CPU as failover: e.g., if GPUs are all down or busy, as a last resort do inference on CPU (terribly slow, but maybe acceptable for a very rare scenario to return something rather than nothing). - Lowering context if memory low – some systems might degrade gracefully by truncating context or switching to a smaller model if resource constrained (not common, but possible strategy).

**Needs:** - More graceful **error recovery**. Perhaps an ability to seamlessly retry a request on another replica mid-stream (almost impossible for long gen because of state). Or replication of KV cache to a backup in real-time (costly, but maybe within a node across GPUs? Unusual). - **Better use of cheap instances:** e.g., use spot instances to carry "overflow" load, but handle when they're yanked (perhaps by keeping a mix of on-demand and spot and always duplicating some requests on both, accepting first answer – probably too expensive to justify).

**Product angle:** - "HA LLM Serving" as a selling point: a service that guarantees X nines of availability by clever redundancy. Could be a niche differentiator (for say, healthcare or mission-critical users who need reliability). - Also a market for **monitoring tools** that can predict failures (like detect a GPU is having memory errors and proactively remove it). - Notably, many end-user apps can tolerate occasional failure (just retry the query), so some companies might not invest heavily here. But enterprise SLAs might demand it.

## 3.6 Cost Engineering

**What it is:** Tracking and optimizing the cost per inference. For cloud, that's GPU-hours; on-prem, electricity and amortization.

**How teams do it:** - Logging token usage per request and multiplying by a cost factor (like \$ per 1000 tokens for that model/GPU). - Attribution: if multiple teams use the cluster, they apportion costs by usage metrics. Tools like **Azure's cost management or custom logging** are used. Some have built internal dashboards of \$ per day per app. - Optimization: once costs are known, they try strategies like

using cheaper hardware (e.g., NVIDIA L4 GPUs are cheaper but slower; perhaps good for smaller models), or adjusting model (quantize to fit more per GPU).

**Pain points:** - Lack of clarity of how much a particular feature usage costs until the bill comes. This can be dangerous if usage skyrockets (some anecdotal cases where an internal LLM tool racked up huge compute bills before anyone noticed). - Hard to map cost to user-level metrics. E.g., if one user of an app makes very long queries, they drive cost more – product might need to charge them accordingly or limit usage, but only if metrics collected support that (like cost per user session). - Tools: As we saw, a number of startups and open tools focus on **observability with cost** <sup>28</sup> <sup>29</sup>. So this is recognized.

**Tricks being used:** - Setting “cost budgets” – e.g., disabling some features if cost this month approaches limit. - Using spot/preemptible GPUs for non-critical inference to cut cost ~70% (with risk of interruptions). - Running models at lower precision or on CPU overnight when responses can be slower but cheaper (rare, but maybe for batch jobs).

**Opportunity:** - Strong demand for *cost dashboards and optimizers*. Possibly integrated with autopilot systems: e.g., a service that automatically switches your model to a smaller one if quality allows for certain queries, to save cost, or autoscaling logic that targets a certain \$/token figure. - Also **billing integration**: If you provide an LLM API to customers, you might want to bill them based on usage – a system that measures and outputs billable metrics (like how OpenAI charges per token) could be offered for those building internal APIs.

In summary, cluster-level issues revolve around **making efficient use of expensive hardware while meeting demand and reliability targets**. Many teams currently solve these in ad-hoc ways (over-provisioning, manual rules, etc.). This suggests room for more systematic solutions – from smarter schedulers to better monitoring and cost management tools – which a startup could productize. The complexity and rapid evolution also mean many organizations will prefer to **buy or use an open platform** rather than maintain deep expertise in this – a promising sign for products in this space.

---

## 4. Startup/Business Opportunity Map

Bringing together the above technical analysis, we can identify clusters of problems that are particularly painful and ripe for product or service solutions. We map these problem clusters, assess who experiences the pain, how it's dealt with today, and propose concrete product ideas. We'll also highlight the strongest opportunities.

### 4.1 Problem Cluster: “Long Context & Memory Management”

- **Pain Intensity:** High. Many teams want to use long documents or chat histories with LLMs, but run into severe performance issues or costs. Memory blows up, latency suffers, or they simply can't enable >4k tokens context due to infra limits.
- **Who feels it:** AI startups offering document analysis, enterprises doing research assistant on long reports, anyone trying to differentiate with “our AI reads all your data not just summary.” Also model providers who want to offer 16k/32k contexts (OpenAI, Anthropic, etc., solved some of it internally at great expense).
- **Current Solutions:** Ad-hoc. Some accept slower speeds or buy more GPUs. Others attempt retrieval as an alternative (which requires changing approach from pure long context to RAG). A

few use vLLM or similar if they know about it. But an enterprise with closed model might not have vLLM support readily, so they might brute force with big memory GPUs.

- **Why a business opportunity:** It's a **clear, recurring technical barrier**. A solution that allows, say, a 7B model to handle 100k tokens with decent speed would enable new use cases. People would pay for that either in software form or as a service (since it saves them from either manually chunking documents or buying 5x more GPUs).

- **Product Ideas:**

- **Long-Context Optimization SDK:** A software library that can be added to existing model servers (or a custom server) which implements efficient paging of KV, maybe compression, and provides APIs to manage context (e.g., load/unload parts of context). It could also include tools to help split documents and use retrieval when context too large.
  - *Who buys:* Enterprises running in-house models wanting to extend context length easily. Possibly open-source with paid enterprise support.
  - *Integration:* Works with popular frameworks (PyTorch, HF Transformers) by hooking the attention kernels, or as a wrap around the model forward pass.
  - *Defensibility:* If it's truly hard to implement and you've got patents or deep know-how (like the authors of PagedAttention), you could lead. Otherwise open-source can catch up (vLLM is open, though not trivial to reimplement).

- **Managed Long-Context Serving:** A cloud service or appliance that clients send their model or data to, which will serve an API with extended context support. For example, upload your model weights, and the service gives you an endpoint / completions call that supports, say, 50k tokens context with only 2x latency of 4k context. Internally it might use a combination of hardware (maybe lots of CPU memory offloading or distributed attention) and software (custom kernels).
  - *Who buys:* Companies that don't mind using a third-party for hosting if it solves a big headache (likely smaller firms or those who can't invest in their own infra). Or sold as an on-prem appliance to sensitive clients.
  - *Defensibility:* Owning a unique solution (maybe low-level kernel IP). And providing it as a service is defensible by operational complexity.

- **Market growing?** Yes, as context lengths become a competitive point (Anthropic moved to 100k, others will follow). Any tool to enable that for more people has growing demand.

**Strongness:** This cluster is strong (clear pain, willing payers) but also competitive – big players are working on it, and open solutions exist. A startup must have either superior tech (e.g., a breakthrough to handle 100k with minimal GPU memory via new algorithms) or more accessible packaging (not everyone can switch to vLLM easily, maybe you make it plug-and-play). Overall business potential: **High** if you have a lead, as it addresses both cost savings and unlocking new use cases.

## 4.2 Problem Cluster: “Multi-Tenant Scheduling & QoS”

- **Pain Intensity:** Medium-High in providers, Medium in enterprises. For anyone offering an AI service to multiple clients or departments, ensuring fairness and meeting different SLAs is a headache. Internally, if one team's job overloads the cluster, others suffer – a political/operational problem.
- **Who feels it:** Cloud API providers (OpenAI, Azure OpenAI, etc.) deal with this at large scale. Also large enterprises where one AI platform is used by multiple internal products. Smaller startups maybe not yet, but if they have tiered customers (free vs paid) they do want to prioritize.
- **Current Solutions:** Over-provision for worst-case to avoid conflict. Some manual segmentation: e.g., dedicate certain GPUs to important tasks, separate out dev/test workloads from production. Using simple queuing and not truly optimizing. Possibly using existing cluster schedulers (Kubernetes) which are not LLM-aware, leading to either underutilization or occasional contention.

- **Why business opportunity:** It's classic "cloud operating system" stuff – building the brain that schedules efficiently and fairly. Given LLM serving is new, the existing general schedulers aren't optimized for it. A product here could save money (by higher utilization instead of siloing hardware) and ensure reliability (prevent noisy neighbors). Particularly enterprises that want to serve many models to many users on shared infrastructure could use something that enforces quotas and priorities elegantly.
- **Product Ideas:**
- **LLM Serving Orchestrator:** Think of it as a specialized scheduler that sits on top of GPU cluster (maybe integrated with Kubernetes or Ray) focusing on LLM workloads. It would have features like per-tenant rate limits, scheduling policies (FCFS, priority levels), dynamic batching across tenants while respecting policies, etc. It might come with a dashboard to configure SLAs ("Tenant A gets 20% of capacity guaranteed", "latency for interactive requests gets P95<2s by limiting batch sizes for them", etc.). Essentially a control plane for multi-LLM, multi-tenant serving.
  - *Who buys:* Enterprises with an AI platform team, AI SaaS providers (who manage many customer models).
  - *Integration:* It could hook into existing infra, e.g., replace Kubernetes scheduler for GPU jobs, or run as a sidecar that decides what requests each running model process should take next.
  - *Defensibility:* If you accumulate data and optimize scheduling policies specific to LLM patterns, that can become a moat (like how Datadog has lots of data on logs to refine their product). Also building trust with enterprise as a reliable piece of infra (which is non-trivial, so established credibility is a moat).
- **Inference API Gateway (with QoS):** A gateway that front-ends all model requests, does admission control, routing to appropriate model instances, and enforces per-customer quotas, even token-based usage limits, etc. It's like an API management tool but tailored to LLM, including awareness of context lengths, streaming, etc. Could be open-source with enterprise features or SaaS.
  - *Who buys:* Companies offering AI APIs or internally exposing models as service to multiple teams.
  - *Integration:* Sits in front of model servers (like an envoy proxy but LLM-smart).
  - *Value-add:* Could automatically produce usage reports (so doubles as billing for cost recoup).
- **5-7 strongest theses:** This cluster is strong if you target the right segment – e.g., an enterprise who doesn't have the scale of OpenAI to build it themselves but has enough internal load to need it (maybe a bank's AI platform for various departments?). Selling them an "AI inference control plane" could be lucrative (they pay for software that saves their engineers time and avoids outages).
- This might be a slightly smaller market than long context or cost-savings because it's relevant once you reach scale. But if adoption of LLMs grows, many medium orgs will hit these multi-user issues.

## 4.3 Problem Cluster: "Observability & Cost Accountability"

- **Pain Intensity:** Medium but rising. Initially people just want it to work; as usage and bills rise, they care a lot about understanding cost drivers and performance issues. Observability is one of those things everyone eventually needs.
- **Who feels it:** Practically any team running non-trivial LLM workloads in production beyond a prototype. Particularly those with tight budgets or latency SLAs. Also finance teams in enterprises wanting to allocate costs to departments using the AI.
- **Current Solutions:** Patchwork – logs with timestamps and token counts, maybe Prometheus metrics for GPU utilization, and a lot of manual analysis. Some might use general APM tools but

they don't have built-in LLM token metrics. New tools (like from TrueFoundry, etc.) are emerging but market still nascent.

• **Why opportunity:** This looks like a classic software tooling market: the Datadog/NewRelic of AI. Those who move first here can become the go-to solution for ML infra monitoring. Given the complexity of LLMs, specialized solutions (monitoring not just if server is up, but how model quality possibly drifts, etc.) can provide high value. People will pay to avoid surprises like a sudden doubling of latency or cost.

• **Product Ideas:**

• **LLM Observatory Platform:** A SaaS (or on-prem) that plugs into your inference servers to collect detailed metrics: per-request and per-token latency, token counts, model invoked, GPU usage timeline, cache hit rates, etc. It then visualizes and alerts on anomalies (e.g., "Latency P99 spiked above SLA at 3PM for Model X", "User Y used 3x more tokens than usual – potential abuse or heavy usage"). It can also compute cost metrics if you input your GPU pricing or use cloud billing APIs. Possibly even tie in generation quality metrics (like scoring outputs or tracking user feedback).

- *Who buys:* ML Ops teams, platform teams – basically anyone deploying LLMs at moderate scale. This could be sold to smaller startups as well who want to monitor their one model, to large companies with many models.
- *Differentiation:* Understanding the LLM-specific metrics (not just treating it as generic web service). For example, correlating prompt length with latency automatically, or showing breakdown of time spent in queue vs compute vs network.
- *Competition:* Early startups and open source are in this space, but no dominant player yet. A good UX and integration with dev workflow (e.g., log traces showing prompt excerpts for slow requests) could win mindshare.

• **Cost Management Service:** A specialized service or tool focusing on cost optimization. It might analyze your usage patterns and recommend actions: "You can save 30% by switching instance type or by quantizing model with minimal accuracy loss" – basically combining observability with optimization suggestions. It could even integrate to automatically implement some optimizations (like auto-scale down unused models to cheaper storage).

- *Who buys:* Companies with big cloud bills for AI – likely mid to large enterprises.
- *This could even be a consulting offering:* analyze an org's AI workload and propose a cost-saving plan (with product tie-in to implement).

• **Strong thesis?:** Observability is almost certainly needed, but it can be hard to monetize as a startup unless you become the one everyone trusts (which is possible, many APM companies exist). The good part is it's not something OpenAI or NVIDIA likely provide (they do internally but not as product), so startups have space.

• This cluster's potential is **High** because it's analogous to DevOps tools which are a huge market. The key is proving value quickly (like show how using the tool prevented an outage or saved X dollars). Many likely happy to pay on a SaaS model per GPU or per million tokens monitored.

## 4.4 Problem Cluster: “Quantization & Model Optimization Tools”

• **Pain Intensity:** Medium. It's technically challenging but the benefits are clear (cost reduction). Many smaller teams might not attempt it yet, but as they scale, they will. Larger tech companies have internal teams for this (so less likely to buy external).

• **Who feels it:** Smaller companies or non-ML specialized enterprises that have a custom model and see inference is costly or slow, but don't have internal research to optimize it. For example, a mid-size company that fine-tuned a Llama-30B for their domain, now wants to deploy on edge devices or just reduce cloud usage.

• **Current Solutions:** Some open-source (like GPTQ scripts) which require technical knowledge. Hugging Face's `optimum` library tries to integrate quantization and compilers, but not always

straightforward. NVIDIA's tools cover their hardware, but e.g., someone using CPU or mobile may struggle.

- **Why opportunity:** If you can make model optimization accessible (like how early web made building websites accessible), you open up more use of these models in various environments (on-phone, in-browser, etc.). And companies will pay for faster inference if it saves them cost or enables new edge deployment (like a bank wanting an offline mobile AI app).
- **Product Ideas:**
- **Automated Model Optimization Service:** Provide a service where you upload a model (or select a known architecture) and it outputs an optimized package: could include quantized weights, compiled runtime (like TensorRT engine or ONNX model), and even benchmark numbers on target hardware. Essentially “press button, get a faster model.” Optionally fine-tune or calibrate if needed as part of process.
  - *Who buys:* Those with custom models to deploy widely, or those who want to support multiple hardware targets without in-house effort (e.g., a software vendor integrating an LLM into their app, needing PC, Mac, mobile support – they could use this service to get optimized binaries for each).
  - *Business model:* Could be one-time service per model or subscription for continuous support (if model updates, you optimize again).
  - *Challenges:* Model IP concerns (some might not want to upload model weights to third-party). Could mitigate by on-prem version or confidential computing.
- **Consulting/Toolkit for Quantization:** Possibly not as scalable, but providing expert guidance and bespoke solutions (like designing a quantization scheme for a unique model, or implementing a custom kernel for a specific use case). Some hardware startups do this to get people using their chips (helping optimize models for their hardware). If a startup had unique quantization IP (say a better algorithm that surpasses GPTQ), they might open-source parts and sell enterprise license/support or custom work.
- **Strongness:** The need is solid, but open-source will keep getting better here (community quickly adopts new quant methods). A startup must either be ahead (with a proprietary technique) or much easier to use. Many might just try open tools if they have moderate expertise. The market of companies completely unable to do it themselves but needing it might be smaller or short-lived. That said, if one becomes known as “the optimization people,” they could gain a lot of business (akin to how some companies specialized in optimizing neural nets for mobile and got acquired).
- Business potential: **Medium.** There’s money here, but the window may close as these techniques commoditize. Might be better as part of a larger offering (like included in an inference platform or cloud service). E.g., Amazon offers “Inferentia” which auto-quantizes behind scenes – so a standalone vendor could be outcompeted if clouds integrate it natively.

## 4.5 Problem Cluster: “Managed LLM Infrastructure (All-in-one)”

*(This cluster is more a combination of many issues – essentially the opportunity to provide a fully managed or packaged inference solution that covers deployment, scaling, monitoring, optimizations, etc. It’s like offering “LLM Platform as a Service.”)*

- **Pain Intensity:** High for those who just want to use LLMs but not build all the plumbing. There’s a reason many still use OpenAI API despite cost – ease.
- **Who feels it:** Enterprises that have data or compliance reasons to run their own models but would prefer not to invest in deep infra. Also startups who want to focus on model/app quality, not on reinventing serving tech.

- **Current Solutions:** Combination of things like Azure ML, Amazon Bedrock, etc., but some find those incomplete or too tied to provider. Some smaller players (e.g., Banana, Anyscale) try to provide simplified infra. None is clearly dominant for on-prem or flexible environments.
- **Why opportunity:** The field is new enough that an agile startup could build a polished product before big slow-moving enterprise software companies do. Customers will pay for anything that saves their engineers time and ensures reliability.
- **Product Ideas:**
- **LLMops Platform:** A unified software that handles model serving (with optimization, multi-gpu), data pipelines (for retrieval, etc.), monitoring, and perhaps even training fine-tunes. Basically the “Snowflake of LLM deployment” or “Databricks for LLM inference.” Many are attempting this; whoever executes well can capture market as usage grows.
- **Specialized Cloud for LLMs:** Not just generic AWS, but a cloud that offers nodes pre-loaded with popular models, automatically scales, with pricing maybe per token or per second rather than per VM. Similar to how some managed database services specialized beyond raw VMs. This is capital intensive but if LLM adoption continues, could carve out share (there are already some AI-focused clouds like CoreWeave focusing on GPUs, but maybe more high-level service).
- **Strongness:** This is broad and competitive, but the market is potentially huge (every company using AI might use such a platform). It combines solutions of all clusters above. The downside is going head-to-head with major cloud providers eventually. But focusing on being multi-cloud or on-prem friendly is a way to differentiate (some enterprises want to run in their own environment, not be locked to a cloud’s offering).
- Many investors are bullish on “dev tools for AI” which this essentially is. So likely multiple entrants, but also possibly multiple winners if each captures different segment.

## 4.6 Problem Cluster: “Domain-Specific & Edge Inference”

- **Pain Intensity:** Medium now, potentially high soon. This covers use cases where you want to run LLMs in constrained or specialized environments – e.g., on IoT devices, in browsers, or with strict safety controls (like healthcare with certain policies).
- **Who feels it:** Companies wanting on-device AI for privacy (e.g., phone manufacturers), or industries with air-gapped networks (defense, some healthcare) who can’t call an API and must run onsite on limited hardware.
- **Current Solutions:** Right now mostly using smaller models or heavily quantized ones (like 4-bit models via llama.cpp on a laptop). Or waiting on big players (Apple, etc., to integrate).
- **Why opportunity:** This is like a new frontier – lots of interest in “LLM at the edge” but few ready solutions for anything beyond toy 7B models. If a product can reliably run a 13B or 30B model on a smartphone or a single board computer with good performance, entire new applications open up. People will pay for that capability (think offline ChatGPT devices, or enterprise appliances for remote locations).
- **Product Ideas:**
- **Edge LLM SDK/Hardware:** Possibly a hardware-software combo specialized for LLM inference in edge scenarios (could be an FPGA or ASIC plus software, but that’s heavy startup lift). Or simpler, an SDK that is highly optimized for available edge hardware (GPU, ARM Neon, etc.), making it easy for developers to plug in a model and run locally with acceleration.
- **Niche model serving:** customizing inference solutions for specific verticals – e.g., a medically certified model server that ensures certain safety checks (like filtering any advice against guidelines) – something an ordinary server won’t do out of the box. Those in regulated fields might value a product that bakes compliance (like logging every output for audit, ensuring no PHI leaves, etc.).

- **Market:** Edge LLM is early but potentially huge (imagine every phone running advanced AI, akin to how mobile app explosion happened when GPUs on phones improved). A startup might create tools now to be the TensorFlow Lite of LLMs.
- This cluster's alignment with business is a bit speculative, but *someone* will solve running bigger models in constrained env, could be open source or chip companies. A nimble startup could fill gaps (like MLC-LLM is an academic project hitting this area; a startup could commercialize such technology).
- Business potential: **Medium (long-term High)** – short term, edge deployments are limited by model capability, but long term as smaller models get better or hardware improves, it will be big. So investing now could pay off if you survive until then.

## 4.7 Strongest Opportunities

From above, **the 5-7 strongest theses** for a business right now appear to be:

- **Long-Context Optimization (High pain, clear \$\$ saving, can be packaged software):** Clear willingness to pay to unlock new use and reduce cost of current use. High technical barrier gives an opening.
- **Inference Orchestration & QoS (Especially for enterprises building internal AI platforms):** These customers have budget and urgent need to manage growing usage. Not many off-the-shelf solutions yet.
- **Observability/Monitoring (AI-specific APM):** A horizontal need across all adopters, could capture a broad market if executed well. Known business models (SaaS) and analogies exist.
- **Managed LLM Ops Platform (Holistic solution):** If a startup can roll many solutions into one and become a one-stop platform, it could be huge (though competition and scope make it challenging). But capturing users early (like how Snowflake captured cloud data warehousing early) can yield big returns.
- **Edge LLM Deployment:** More future oriented, but a risky bet that could pay off massively if AR/VR, mobile, IoT demands on-device AI. If a startup becomes leader here, large companies (chip makers, etc.) might acquire for their expertise.
- **Model Optimization & Cost Reduction Tools:** There's money to be made saving money for others. However, this might be a shorter-term opportunity until mainstream frameworks incorporate these optimizations. Still, in next 1-2 years, many would pay for 2x cost reduction via better quantization or compilers.
- **Specialized Cloud Services (like a latency-optimized inference cloud):** Possibly strong if differentiated enough (e.g., guarantee <50ms response for certain models via custom hardware etc.). It is capital heavy though, so maybe not for a small startup without major backing.

In prioritizing, the **top 3** might be: 1. **Observability/Cost platform** – because it's software-only, touches every deployment, and relatively less head-on competition yet. 2. **Long context & memory solutions** – technical moat possible, many willing customers with immediate pain. 3. **QoS orchestration** – because enterprises will need it, and if you integrate into their pipeline early, it's sticky.

These combine strong pain + growing market + not fully solved by giants (yet). A strategy for a startup could be to tackle one as a beachhead (e.g., release an open-source long-context server that gains community traction), then build complementary features (monitoring, scheduling, etc.) around it as a paid offering, effectively evolving into the all-in-one platform with a user base already in place.

---

# 5. Annotated Bibliography

Below we list key sources (papers, blogs, repos, talks) that informed this report, with brief summaries and relevance:

## Papers

- **Kwon et al. (2023) – *Efficient Memory Management for LLM Serving with PagedAttention*** 88 89: Introduces **PagedAttention** and vLLM. Shows that conventional KV cache handling wastes 60–80% memory due to fragmentation, and how paging + sharing can improve throughput 2-4× without latency loss 88 89. *Research-grade but implemented in vLLM (SOSP 2023), forms the foundation for modern high-throughput inference (section 2.1, 2.6).*
- **Joshi et al. (2025) – *PagedAttention meets FlexAttention*** 1 40: An IBM-Columbia work integrating paged KV management with a fused attention kernel via PyTorch 2.x’s FlexAttn. Confirms fragmentation issues and reports near-linear latency scaling up to 2k tokens with global cache 58 1. Indicates industry’s interest in long-context optimization (sections 2.1, 4.1).
- **Gaur (2024, Medium) – *Understanding ORCA*** 66 24: Summary of ORCA (a Microsoft Research inference system). Explains **iteration-level scheduling** to solve batch inefficiencies 66 24 and introduces **selective batching** to combine different sequence lengths 67 68. *Practice-grade concept implemented in some form in vLLM; supports sections 2.2 and 2.6 on scheduling.*
- **Li et al. (2022) – *Contrastive Search Is What You Need*** 90 (arXiv): Proposed **contrastive decoding**, using a large “expert” and small “amateur” model to choose tokens that maximize the difference in prediction vs a constraint 91 79. Shows improved output quality (less repetition). *Research-grade decoding method, illustrating the trend of multi-model decoding algorithms (section 2.4).*
- **Dettmers et al. (2022) – *GPTQ: Accurate Post-Training Quantization*** (arXiv): Developed GPTQ algorithm for 4-bit quantization of transformers. Demonstrated minimal loss on GPT models. Paved the way for broad community use of 4-bit models. *Supports section 2.3 on quantization.* (No direct excerpt above, but known reference).
- **AWQ (Lin et al. 2023, MLSys)** 14 : Activation-aware Weight Quantization method allowing 4-bit weight quant with negligible drop, by addressing outlier channels 14. *Supports quantization advances mentioned in section 2.3.*
- **Microsoft (2023) – *Inference Techniques like Speculative Decoding*** 92 16 : A blog referencing that pairing a smaller model to guess ahead can 2-3× speed up generation 92 16. *Validates speculative decoding performance (section 2.4).*
- **Shah et al. (2023) – *LLMVisor: Latency Attribution for Multi-Tenant LLMs*** 26 (OpenReview): Proposes a tool to attribute GPU time per tenant and enforce fairness in scheduling 26. *Shows industry recognition of multi-tenant fairness issue (section 2.2, 3.2).*

## Engineering Blogs & Articles

- **vLLM Team Blog (2023) – Easy, Fast, Cheap LLM Serving** 8 11 : Describes vLLM's benefits. Notably, shows *24x throughput vs HuggingFace* and *3x vs TGI* 93 7 , and <4% memory waste with PagedAttention vs 60–80% before 8 11 . Mentions prefix sharing improving parallel sampling throughput 2.2x 57 . Supports multiple points: *huge throughput gains (section 1.3, 2.1), memory waste stats (1.3, 2.1), and prefix cache utility (2.6)*.
- **Aleksa Gordić (2025) – Inside vLLM** 94 43 : Deep dive on vLLM internals. Highlights scheduler options (FCFS vs priority) and the block-based KV cache manager with free\_block\_queue for pages 43 . Informs sections 2.2 (scheduling) and 2.6 (cache design).
- **Medium (Sep 2025) – vLLM vs TGI vs TensorRT-LLM (Bhagya Rana)** 95 96 : A comparative intro that succinctly says *vLLM = memory-savvy speed (PagedAttention + continuous batching, prefix caching, speculative decoding)* 95 ; *TensorRT-LLM = lowest latency if willing to compile on NVIDIA; TGI = easy HF integration. Used in section 1.3 and exec summary to characterize runtime trade-offs.* (Note: content was paywalled; we gleaned summary from snippet).
- **Hivenet Blog (2023) – vLLM vs others** 97 98 : Discusses why teams pick TensorRT-LLM: highest perf on NVIDIA, but requires engine building; notes it also benefits from advanced attention like PagedAttention 97 . Also mentions **MLC-LLM** as emerging but with limitations (compilation need, etc.) 99 . Used in section 1.3 (runtime layer) to show each engine's niche and that even TensorRT is adopting similar KV management ideas.
- **Shashank Guda (2023) – LLM Deployments Aren't Plug & Play** 18 50 : An overview of enterprise challenges. Provides concrete figures: *Llama 3.1 405B requires 800+ GB GPU memory (dozens of GPUs)* 18 ; *cost of an 8xA100 server \$33/hour (~\$287k/year)* 100 . Also shows *throughput vs latency curve: as concurrency increases, throughput increases but latency also climbs, thus need to find optimal point and use techniques like dynamic batching; vLLM's continuous batching gave up to 23x throughput and reduced median latency vs naive* 50 5 . These support the need for batching (2.2) and illustrate cost magnitude (3.1) and scalability issues (3.1, 3.2).
- **NVIDIA Technical Blog (Sep 2025) – Run:AI & Dynamo for Multi-Node LLM Inference** 101 72 : Announces integration of NVIDIA Dynamo (inference framework) with Run:AI scheduler. States features: *disaggregated prefill/decode, dynamic GPU scheduling, LLM-aware routing (preventing redundant recompute), KV cache offloading* 101 . Emphasizes need for gang scheduling and topology-aware placement in multi-node to avoid idle partial deployments and minimize cross-node latency 102 103 . Used in section 2.2 (mentioning disaggregated stages) and section 3.4/3.2 (multi-node scheduling and topology). Shows state-of-art industry solution for cluster scheduling (informs opportunity mapping as well).
- **HuggingFace Blog (Jul 2023) – Generating Human-Level Text with Contrastive Search** 104 91 : Explains contrastive search as balancing model's fluency vs diversity by avoiding model's highest-confidence (possibly repetitive) tokens unless they are much more likely than alternatives 91 . Though about decoding quality, it's an example of new decoding algorithms (section 2.4).
- **Reddit (2023) – Which is faster: vLLM, TGI, or TensorRT?** 105 35 : Community discussion indicating *all three are similar at moderate load, with TGI a bit faster at low QPS, and vLLM fastest at high concurrency; TensorRT-LLM lowest latency for batch-1 on supported GPUs* 35 . MarkTechPost

summary confirms *TensorRT-LLM on H100 can get <10ms first token in batch-1, but lower throughput at scale* <sup>39</sup>. Corroborates engine trade-offs (sec 1.3, Exec summary bullet).

## Repos & Issues

- **vLLM GitHub** <sup>37</sup> <sup>106</sup>: The README highlights “*efficient management of attention KV memory with PagedAttention; continuous batching; fast execution with CUDA graphs*” <sup>37</sup>. Also news of v1 adding “*zero-overhead prefix caching*” <sup>85</sup>. Used to ensure accuracy on vLLM features (prefix caching – sec 2.6) and to note use of CUDA graphs (sec 1.2).
- **state-spaces/mamba GitHub** (not excerpted above): Provided insight that Mamba is being practically experimented with, indicating SSMs potential. Their README likely notes throughput improvements and linear scaling. *Relevant to sec 2.3 architecture discussion (we cited arXiv instead)*.
- **NVIDIA/TensorRT-LLM GitHub** (not directly quoted): Confirms it’s open-source now, focusing on optimizations for LLM (supports sections 1.3 and 2.2/2.3 by implying broad adoption of such engines).
- **Huggingface TGI Issues** (not quoted): Many issues discuss multi-client usage, memory growth, etc., implying the need for better multi-tenant handling. *This background informed sec 2.6 and 3.2 but not directly cited due to time.*

## Talks / Videos

- **Emergent Mind (2023) – LLM Inference Scheduling Overview** <sup>107</sup> <sup>108</sup>: An article on EmergentMind (possibly by swyx) that outlines scheduling framework. It likely refers to ORCA and others. We saw a snippet referencing “*LLM inference scheduling allocates resources and batches requests to manage memory constraints and variable output...*” <sup>107</sup>. *Helps in sec 2.2 (as evidence that inference scheduling is being recognized in discourse)*.
- **YouTube: Speculative Decoding explained** (not excerpted): Provided general understanding of how the draft model and main model interplay, which we used in sec 2.4.
- **AI Engineering Summit talks (2023)** (not specific excerpt): Many talks (by Meta, OpenAI, etc.) on serving at scale influenced our understanding of current best practices and problems (like comments publicly made by OpenAI on speculative decoding use, by Meta on Llama scaling). *No direct citation, but background knowledge supporting multiple sections.*

This bibliography shows how we pulled from both scholarly works (for validated techniques and data points) and real-world engineering discussions (for practical pains and solutions). Each source has been tied into sections of this report, providing evidence for claims like “memory waste 60–80%” <sup>8</sup>, “vLLM 24x throughput vs HF” <sup>93</sup>, “throughput vs latency trade-off” <sup>50</sup>, and many more as cited inline. Together, they map the frontier of LLM inference as of late 2025.

---

# 6. Glossary & Concept Graph

## Glossary of Key Terms

- **LLM (Large Language Model):** A neural network model with billions of parameters trained on vast text data, capable of generating human-like text. E.g., GPT-3, LLaMA, etc.
- **Inference:** The process of using a trained model to generate outputs (as opposed to training). In LLM context, it often means generating text from a prompt.
- **Throughput:** Amount of work done per time (e.g., tokens generated per second). Critical metric for serving many requests.
- **Latency:** Delay to get a result. For LLM, often measured as time to first token or time to complete response.
- **KV Cache (Key-Value Cache):** In transformer models, the stored key and value vectors from previous tokens, used to speed up attention calculations for next token.
- **PagedAttention:** An approach that breaks the KV cache into fixed-size blocks (“pages”) allowing flexible memory management (as if GPU memory had paging like an OS) <sup>54</sup>.
- **Continuous Batching:** Dynamically combining incoming requests or token computations into batches on the fly, rather than fixed batches. Keeps GPUs busy with multiple tasks interleaved.
- **Speculative Decoding:** Using a fast “draft” model to predict a few tokens ahead, and having the large model verify them, thus skipping some large model forward passes if the draft is correct <sup>16</sup>.
- **Quantization:** Reducing numerical precision of model weights/activations (e.g., 16-bit to 8-bit or 4-bit) to decrease memory and compute usage <sup>14</sup>. Can reduce model accuracy if not done carefully.
- **Tensor Parallelism:** Splitting a model’s layers (especially weight matrices) across multiple GPUs, each computing part of the output. Requires GPUs to communicate at every layer (all-reduce or similar).
- **Pipeline Parallelism:** Splitting the sequence of layers across GPUs. The input flows through GPU1’s layers, then GPU2’s, etc. Allows fitting bigger model, but introduces pipeline latency.
- **NVLink / NVSwitch:** High-bandwidth interconnects between GPUs (NVLink typically connects 2 GPUs ~50–100 GB/s, NVSwitch is like an all-to-all network inside a server ~600+ GB/s). Much faster than PCIe (~16 GB/s).
- **Batching:** Grouping multiple inputs together to process in one forward pass (leveraging parallelism). In LLM serving, often refers to grouping multiple requests or tokens.
- **Mixture-of-Experts (MoE):** A model architecture with multiple expert sub-models where only a few are active for each input. Reduces computation per token by not running all parameters every time, but adds routing overhead.
- **FlashAttention:** An optimized attention kernel that uses tiling and recomputation to reduce memory usage and increase speed, especially for long sequences <sup>109</sup>.
- **Beam Search:** A decoding algorithm that keeps multiple candidate sequences at each step and eventually chooses the highest probability one. Ensures higher-likelihood outputs but can be slower and less diverse.
- **Nucleus/Top-p Sampling:** A decoding method that at each step restricts sampling to the smallest set of tokens whose cumulative probability exceeds p (like 0.9). Aimed at balancing coherence and randomness.
- **Multitenancy:** Serving multiple independent users or clients on the same system. Implies need for isolation (one user’s heavy usage shouldn’t degrade others, data separation, etc.).
- **Autoscaling:** Automatically adjusting the number of running instances (servers/GPUs) based on load, to handle demand while minimizing idle resources.

- **SLA (Service Level Agreement):** Target reliability or performance metrics (like 99th percentile latency) that a service promises to meet. Drives engineering decisions on capacity and scheduling.
- **CUDA Graphs:** A feature in NVIDIA CUDA that allows capturing a sequence of GPU operations and replaying it with less overhead. Useful for eliminating per-iteration launch costs in steady loops (like generation).

## Concept Graph (Relationships)

- **Model Architecture vs Hardware:** The size and structure of the model (e.g., 400B parameters vs 7B, or using MoE vs dense) determine memory and compute needs. Hardware constraints (GPU memory, compute throughput, network speed) in turn limit how big a model can be served on one device or require splitting (hence influencing architecture choices like needing MoE to reduce per-GPU load).
- **Model (and context length) → Runtime Techniques:** A model with long context (say 32k) pushes the runtime layer to implement things like PagedAttention to manage memory, and perhaps flash-attention kernels to handle compute. A smaller context model wouldn't necessitate those.
- **Runtime ↔ Hardware:** The serving software (runtime) tries to maximize hardware usage via batching, parallelism, etc. If hardware has NVLink, runtime can do tensor parallel across GPUs; if not, runtime might avoid splitting and instead run separate models per GPU. Hardware features like FP8 dictate runtime enabling mixed precision for speed.
- **Runtime ↔ Infrastructure:** The inference server (runtime) typically runs on a node; the cluster infrastructure (e.g., Kubernetes, or a scheduler) decides how many instances, where to place them, etc. If runtime can batch across requests, the cluster might need fewer total instances (because each handles more throughput). Conversely, if infrastructure auto-scales out instances, each runtime might use smaller batch sizes. They influence each other's efficacy – e.g., if autoscaler frequently turns off instances, runtime might not reach steady high throughput.
- **Infrastructure ↔ Workload Patterns:** If the workload is interactive chat (bursty, latency-sensitive), the infrastructure might allocate more headroom (extra GPUs) to ensure low latency and implement priority scheduling. If workload is batch (steady, throughput-focused), infrastructure can run at higher utilization and queue requests. So the type of product usage dictates infrastructure tuning (SLA vs cost trade-off).
- **Product Requirements ↔ Model and Safety:** Product constraints (like needing real-time answers, or ensuring no disallowed content) can affect everything down the stack: e.g., a strict latency SLA might force using a smaller model (model choice) or quantization (model optimization) to meet timing. Safety requirements might require an extra filtering model to run alongside (impacting runtime throughput and needing more hardware).
- **Business Constraints ↔ Infrastructure & Optimizations:** If the business needs to cut costs, that flows down to enabling quantization, reducing redundancy, and better scheduling to use fewer GPUs (sections 2.3, 2.2). If the business priority is quality or new features, they might allocate more hardware to run larger models or longer contexts, which then requires the technical solutions (like memory management) to do so efficiently.

In summary, the relationships can be seen as a flow:

**User/Product demands (e.g., fast chat for many users, or high accuracy, or privacy)** ⇒ influence **Model choice (size, architecture) and features (context length)** ⇒ which impose needs on **Runtime/Serving (for performance, memory, scheduling)** ⇒ which stress **Hardware (GPU memory, compute)** in certain ways ⇒ requiring changes in **infrastructure deployment (more GPUs, distributed, specific hardware)**. Meanwhile, **business considerations (cost budgets, reliability targets)** feed back to

choose different optimizations at all levels (maybe quantize model, or double infrastructure for redundancy, etc.).

This concept graph shows a highly interconnected system: optimizing LLM inference is a multi-layer puzzle, where a bottleneck at one layer (say memory) can be addressed by innovation in that layer (paging) or by shifting load to another layer (distribute to more GPUs) – each with consequences on cost and complexity. Successful engineering – and thus startup solutions – often means balancing these trade-offs to meet a specific set of product and business goals.

---

1 9 31 32 40 49 58 59 62 Paged Attention Meets FlexAttention: Unlocking Long-Context Efficiency in Deployed Inference

<https://arxiv.org/html/2506.07311v1>

2 6 7 8 10 11 54 55 56 57 60 93 vLLM: Easy, Fast, and Cheap LLM Serving with PagedAttention | vLLM Blog

<https://blog.vllm.ai/2023/06/20/vllm.html>

3 4 22 23 24 25 41 42 63 64 65 66 67 68 69 70 Understanding Orca. As LLMs gained popularity, service... | by Pratishtha Gaur | Medium

<https://medium.com/@pratishtaghagaur03/understanding-orca-a-research-synopsis-037797002039>

5 18 50 51 100 LLM Deployments Aren't Plug & Play | by Shashank Guda | Medium

<https://shashankguda.medium.com/llm-deployments-arent-plug-play-78e383c6561c>

12 75 Demystifying LLM Quantization: GPTQ, AWQ, and GGUF Explained

<https://www.linkedin.com/pulse/demystifying-llm-quantization-gptq-awq-gguf-explained-xiao-fei-zhang-1lmb>

13 Comparison: GPTQ vs AWQ vs SmoothQuant - ApX Machine Learning

<https://apxml.com/courses/practical-llm-quantization/chapter-3-advanced-ptq-techniques/comparing-advanced-ptq>

14 [PDF] AWQ: Activation-aware Weight Quantization for On-Device LLM ...

[https://proceedings.mlsys.org/paper\\_files/paper/2024/file/42a452cbafa9dd64e9ba4aa95cc1ef21-Paper-Conference.pdf](https://proceedings.mlsys.org/paper_files/paper/2024/file/42a452cbafa9dd64e9ba4aa95cc1ef21-Paper-Conference.pdf)

15 [2312.00752] Mamba: Linear-Time Sequence Modeling with Selective State Spaces

<https://arxiv.org/abs/2312.00752>

16 17 How Speculative Decoding Boosts vLLM Performance by up to 2.8x

<https://blog.vllm.ai/2024/10/17/spec-decode.html>

19 20 21 71 72 73 74 83 101 102 103 Smart Multi-Node Scheduling for Fast and Efficient LLM

Inference with NVIDIA Run:ai and NVIDIA Dynamo | NVIDIA Technical Blog

<https://developer.nvidia.com/blog/smarter-multi-node-scheduling-for-fast-and-efficient-llm-inference-with-nvidia-runai-and-nvidia-dynamo/>

26 27 [PDF] A Real-Time Latency Attribution Model for Multi-Tenant LLM Serving

<https://openreview.net/pdf?id=VeAJmE2ZJU>

28 AI Observability: How to Keep LLMs, RAG, and Agents Reliable in ...

<https://www.logicmonitor.com/blog/ai-observability>

29 Monitoring Latency and Cost in LLM Operations: Essential Metrics ...

<https://www.getmaxim.ai/articles/monitoring-latency-and-cost-in-llm-operations-essential-metrics-for-success/>

30 What is LLM Observability ? Complete Guide - TrueFoundry

<https://www.truefoundry.com/blog/what-is-llm-observability>

33 61 FlashAttention-2: Faster Attention with Better Parallelism and Work...

<https://openreview.net/forum?id=mZn2Xyh9Ec>

34 FlashAttention 2: making Transformers 800% faster w/o approximation

<https://www.latent.space/p/flashattention>

35 vLLM vs Ollama vs llama.cpp vs TGI vs TensorRT-LLM: 2025 Guide

<https://itecsonline.com/post/vllm-vs-ollama-vs-llama.cpp-vs-tgi-vs-tensorrt>

36 39 vLLM vs TensorRT-LLM vs HF TGI vs LMDeploy, A Deep Technical ...

<https://www.marktechpost.com/2025/11/19/vllm-vs-tensorrt-llm-vs-hf-tgi-vs-lmdeploy-a-deep-technical-comparison-for-production-llm-inference/>

- 37 85 86 106 GitHub - vllm-project/vllm: A high-throughput and memory-efficient inference and serving engine for LLMs  
<https://github.com/vllm-project/vllm>
- 38 vLLM vs TensorRT-LLM vs HF TGI vs LMDeploy: LLM Inference  
<https://thinktools.ai/blog/vllm-vs-tensorrtllm-vs-hf-tgi-vs-lmdeploy-llm-inference>
- 43 44 94 Inside vLLM: Anatomy of a High-Throughput LLM Inference System - Aleksa Gordić  
<https://www.aleksagordic.com/blog/vllm>
- 45 46 108 [PDF] DeltaZip: Efficient Serving of Multiple Full-Model-Tuned LLMs  
<https://anakli.inf.ethz.ch/papers/deltazip.pdf>
- 47 [2507.11181] Mixture of Experts in Large Language Models - arXiv  
<https://arxiv.org/abs/2507.11181>
- 48 Mixture of Experts (MoE) - Sebastian Raschka  
[https://sebastianraschka.com/llms-from-scratch/ch04/07\\_moe/](https://sebastianraschka.com/llms-from-scratch/ch04/07_moe/)
- 52 53 The Future of LLM Routing: On-device, Edge AI, and Federated ...  
<https://www.requesty.ai/blog/the-future-of-llm-routing-on-device-edge-ai-and-federated-models-1751656903>
- 76 NeurIPS Post Training Quantization of Large Language Models with ...  
<https://nips.cc/virtual/2024/106487>
- 77 Deploy large language models on AWS Inferentia2 using large ...  
<https://aws.amazon.com/blogs/machine-learning/deploy-large-language-models-on-aws-inferentia2-using-large-model-inference-containers/>
- 78 Serving LLMs on AWS EC2 with Inferentia chip, Neuron SDK and ...  
<https://medium.com/techbull/serving-llms-on-aws-ec2-with-inferentia-chip-neuron-sdk-and-dlami-8c4b937f175b>
- 79 [PDF] A Thorough Examination of Decoding Methods in the Era of LLMs  
<https://aclanthology.org/2024.emnlp-main.489.pdf>
- 80 Explaining and improving contrastive decoding by extrapolating the ...  
<https://www.amazon.science/publications/explaining-and-improving-contrastive-decoding-by-extrapolating-the-probabilities-of-a-huge-and-hypothetical-lm>
- 81 82 Speculative Cascades: Unlocking Smarter, Faster LLM Inference  
<https://joshuaberkowitz.us/blog/news-1/speculative-cascades-unlocking-smarter-faster-llm-inference-1107>
- 84 An LLM Service System with Efficient and Effective Multi-Tenant KV ...  
<https://arxiv.org/html/2503.16525v2>
- 87 [PDF] Multi-tenant Machine Learning Model Serving Systems on GPU ...  
<https://digital.lib.washington.edu/bitstreams/f616e12e-d58c-41d3-a3b8-ef66cfbb9c8a/download>
- 88 89 [2309.06180] Efficient Memory Management for Large Language Model Serving with PagedAttention  
<https://arxiv.org/abs/2309.06180>
- 90 Contrastive Search Is What You Need For Neural Text Generation  
<https://arxiv.org/abs/2210.14140>
- 91 Contrastive Decoding Improves Reasoning in Large Language ...  
[https://www.reddit.com/r/LocalLLaMA/comments/16mwcch/contrastive\\_decoding\\_improves\\_reasoning\\_in\\_large/](https://www.reddit.com/r/LocalLLaMA/comments/16mwcch/contrastive_decoding_improves_reasoning_in_large/)
- 92 Blogs - Speculative Decoding - Infocusp Innovations  
<https://www.infocusp.com/blogs/speculative-decoding/>

95 96 vLLM vs TGI vs TensorRT-LLM: Tokens/sec Showdown | by Bhagya Rana | Medium  
<https://medium.com/@bhagyaranana80/vllm-vs-tgi-vs-tensorrt-llm-tokens-sec-showdown-1171a5ed326e>

97 98 99 vLLM vs TGI vs TensorRT-LLM vs Ollama | Hivenet  
<https://compute.hivenet.com/post/vllm-vs-tgi-vs-tensorrt-llm-vs-ollama>

104 Generating Human-level Text with Contrastive Search in Transformers  
<https://huggingface.co/blog/introducing-csearch>

105 TGI vs. TensorRT LLM: The Best Inference Library for Large ...  
<https://www.inferless.com/learn/tgi-vs-tensorrt-llm-the-best-inference-library-for-large-language-models>

107 LLM Inference Scheduling Overview - Emergent Mind  
<https://www.emergentmind.com/topics/llm-inference-scheduling>

109 FlashAttention-2: Faster Attention with Better Parallelism and Work ...  
<https://crfm.stanford.edu/2023/07/17/flash2.html>